

DosDetection

December 9, 2017

1 i523 Big Data Analytics Project

1.0.1 HID224 - Neha Rawat

- 1.1 * Big Data Analytics in Detection of DDoS (Distributed Denial-of-Service) attacks**
- * With the increase in internet traffic, threats on the network have also increased. Denial-of-service attacks are cyber attacks wherein a perpetrator, due to any kind of malicious intent, tries to make a resource on the network unavailable to its intended users and carries it out by swamping the system or resource with excess requests in order to overload it and prevent users from accessing it. A much more dangerous variety of such an attack is if it is distributed i.e. coming from various sources. Big Data analytics, however, can be used to detect such attacks by having the ability to store the voluminous logs of such attacks and using the data and machine learning techniques to design an anomaly detection system (using a classification model) to detect and prevent these attacks. This project will aim to explore such classification models, design and train the most optimum model and display its effects using a DDoS logs dataset.**

1.1.1 Data Description

- We will be using the dataset used for the KDD Cup of 1999, which has been derived from the 1998 DARPA Intrusion Detection Evaluation Program dataset. This was prepared and managed by MIT Lincoln Labs. The objective was to survey and evaluate research in intrusion detection. The data includes a wide variety of intrusions simulated in a military network environment.
- Lincoln Labs set up an environment to acquire nine weeks of raw TCP dump data for a local-area network (LAN) simulating a typical U.S. Air Force LAN. They operated the LAN as if it were a true Air Force environment, but peppered it with multiple attacks. The raw training data was about four gigabytes of compressed binary TCP dump data from seven weeks of network traffic. This was processed into about five million connection records.
- KDD Cup 1999 data link: [link_here](#)
- Original DARPA datasets link: [link_here](#)
- For the data analysis and demonstration of our model, we will use 10 percent of the KDD Cup'99 training dataset (~500k rows) and test it on 10 percent of the test dataset (~300k rows).
- The data consists of five types of network traffic logs: normal traffic, DOS/DDoS - Denial of Service traffic, R2L- unauthorized access from a remote machine, U2R- unauthorized access to local superuser (root) privileges and probing- surveillance and other probing. This makes

the task of analysis and prediction much more realistic, as we attempt to classify the "bad traffic" from the "normal traffic". Another important point to note is that the datasets contain a total of 24 training attack types (subsets of the above four), with an additional 14 types in the test data only. The purpose behind adding the 14 extra types in the test dataset was to prove that most new attacks are just variants of pre-existing attack types and hence can be sufficiently detected using data from the latter.

```
In [1]: # Reading dataset using Pandas
        #Datasets being used for the analysis are the 10 percent datasets placed on Google Drive
        #Place the code and datasets in your default directory for Jupyter Notebook

        #Loading the 10 percent training dataset
import pandas as pd
labels_data = ["duration","protocol_type","service","flag","src_bytes", "dst_bytes","land",
               "logged_in","num_compromised","root_shell","su_attempted","num_root","num_file_created",
               "is_host_login","is_guest_login","count","srv_count","serror_rate","srv_serror_rate",
               "diff_srv_rate","srv_diff_host_rate","dst_host_count","dst_host_srv_count","dst_host_same_srv_rate",
               "dst_host_diff_srv_rate","dst_host_serror_rate","dst_host_srv_serror_rate","dst_host_count_per_ip"]
kdd_train_10percent = pd.read_csv("kddcup_data_10_percent.csv", header=None, names = labels_data)
kdd_train_10percent.head(10)
```

```
Out[1]:
```

	duration	protocol_type	service	flag	src_bytes	dst_bytes	land	\
0	0	tcp	http	SF	181	5450	0	
1	0	tcp	http	SF	239	486	0	
2	0	tcp	http	SF	235	1337	0	
3	0	tcp	http	SF	219	1337	0	
4	0	tcp	http	SF	217	2032	0	
5	0	tcp	http	SF	217	2032	0	
6	0	tcp	http	SF	212	1940	0	
7	0	tcp	http	SF	159	4087	0	
8	0	tcp	http	SF	210	151	0	
9	0	tcp	http	SF	212	786	0	

	wrong_fragment	urgent	hot	...	dst_host_srv_count	\
0	0	0	0	...	9	
1	0	0	0	...	19	
2	0	0	0	...	29	
3	0	0	0	...	39	
4	0	0	0	...	49	
5	0	0	0	...	59	
6	0	0	0	...	69	
7	0	0	0	...	79	
8	0	0	0	...	89	
9	0	0	1	...	99	

	dst_host_same_srv_rate	dst_host_diff_srv_rate	\
0	1.0	0.0	
1	1.0	0.0	

2	1.0	0.0
3	1.0	0.0
4	1.0	0.0
5	1.0	0.0
6	1.0	0.0
7	1.0	0.0
8	1.0	0.0
9	1.0	0.0

	dst_host_same_src_port_rate	dst_host_srv_diff_host_rate \
0	0.11	0.00
1	0.05	0.00
2	0.03	0.00
3	0.03	0.00
4	0.02	0.00
5	0.02	0.00
6	1.00	0.04
7	0.09	0.04
8	0.12	0.04
9	0.12	0.05

	dst_host_serror_rate	dst_host_srv_serror_rate	dst_host_rerror_rate \
0	0.0	0.0	0.0
1	0.0	0.0	0.0
2	0.0	0.0	0.0
3	0.0	0.0	0.0
4	0.0	0.0	0.0
5	0.0	0.0	0.0
6	0.0	0.0	0.0
7	0.0	0.0	0.0
8	0.0	0.0	0.0
9	0.0	0.0	0.0

	dst_host_srv_rerror_rate	attack_type
0	0.0	normal.
1	0.0	normal.
2	0.0	normal.
3	0.0	normal.
4	0.0	normal.
5	0.0	normal.
6	0.0	normal.
7	0.0	normal.
8	0.0	normal.
9	0.0	normal.

[10 rows x 42 columns]

In [2]: kdd_train_10percent.describe()

```

Out[2]:
      duration      src_bytes      dst_bytes      land \
count 494021.000000 4.940210e+05 4.940210e+05 494021.000000
mean   47.979302 3.025610e+03 8.685324e+02 0.000045
std   707.746472 9.882181e+05 3.304000e+04 0.006673
min     0.000000 0.000000e+00 0.000000e+00 0.000000
25%     0.000000 4.500000e+01 0.000000e+00 0.000000
50%     0.000000 5.200000e+02 0.000000e+00 0.000000
75%     0.000000 1.032000e+03 0.000000e+00 0.000000
max    58329.000000 6.933756e+08 5.155468e+06 1.000000

      wrong_fragment      urgent      hot      num_failed_logins \
count 494021.000000 494021.000000 494021.000000 494021.000000
mean   0.006433 0.000014 0.034519 0.000152
std   0.134805 0.005510 0.782103 0.015520
min     0.000000 0.000000 0.000000 0.000000
25%     0.000000 0.000000 0.000000 0.000000
50%     0.000000 0.000000 0.000000 0.000000
75%     0.000000 0.000000 0.000000 0.000000
max     3.000000 3.000000 30.000000 5.000000

      logged_in      num_compromised      ... \
count 494021.000000 494021.000000 ...
mean   0.148247 0.010212 ...
std   0.355345 1.798326 ...
min     0.000000 0.000000 ...
25%     0.000000 0.000000 ...
50%     0.000000 0.000000 ...
75%     0.000000 0.000000 ...
max     1.000000 884.000000 ...

      dst_host_count      dst_host_srv_count      dst_host_same_srv_rate \
count 494021.000000 494021.000000 494021.000000
mean   232.470778 188.665670 0.753780
std   64.745380 106.040437 0.410781
min     0.000000 0.000000 0.000000
25%    255.000000 46.000000 0.410000
50%    255.000000 255.000000 1.000000
75%    255.000000 255.000000 1.000000
max    255.000000 255.000000 1.000000

      dst_host_diff_srv_rate      dst_host_same_src_port_rate \
count 494021.000000 494021.000000
mean   0.030906 0.601935
std   0.109259 0.481309
min     0.000000 0.000000
25%     0.000000 0.000000
50%     0.000000 1.000000
75%     0.040000 1.000000

```

max	1.000000	1.000000
-----	----------	----------

	dst_host_srv_diff_host_rate	dst_host_serror_rate \
count	494021.000000	494021.000000
mean	0.006684	0.176754
std	0.042133	0.380593
min	0.000000	0.000000
25%	0.000000	0.000000
50%	0.000000	0.000000
75%	0.000000	0.000000
max	1.000000	1.000000

	dst_host_srv_serror_rate	dst_host_rerror_rate \
count	494021.000000	494021.000000
mean	0.176443	0.058118
std	0.380919	0.230590
min	0.000000	0.000000
25%	0.000000	0.000000
50%	0.000000	0.000000
75%	0.000000	0.000000
max	1.000000	1.000000

	dst_host_srv_rerror_rate
count	494021.000000
mean	0.057412
std	0.230140
min	0.000000
25%	0.000000
50%	0.000000
75%	0.000000
max	1.000000

[8 rows x 38 columns]

In [2]: *#Loading the 10 percent testing dataset*

```
kdd_test_10percent = pd.read_csv("corrected.csv", header=None, names = labels_data)
kdd_test_10percent.head(10)
```

```
Out[2]:
```

	duration	protocol_type	service	flag	src_bytes	dst_bytes	land	\
0	0	udp	private	SF	105	146	0	
1	0	udp	private	SF	105	146	0	
2	0	udp	private	SF	105	146	0	
3	0	udp	private	SF	105	146	0	
4	0	udp	private	SF	105	146	0	
5	0	udp	private	SF	105	146	0	
6	0	udp	domain_u	SF	29	0	0	
7	0	udp	private	SF	105	146	0	
8	0	udp	private	SF	105	146	0	

9	0	tcp	http	SF	223	185	0
	wrong_fragment	urgent	hot	...	dst_host_srv_count \		
0	0	0	0	...	254		
1	0	0	0	...	254		
2	0	0	0	...	254		
3	0	0	0	...	254		
4	0	0	0	...	254		
5	0	0	0	...	255		
6	0	0	0	...	3		
7	0	0	0	...	253		
8	0	0	0	...	254		
9	0	0	0	...	255		

	dst_host_same_srv_rate	dst_host_diff_srv_rate	\
0	1.00	0.01	
1	1.00	0.01	
2	1.00	0.01	
3	1.00	0.01	
4	1.00	0.01	
5	1.00	0.00	
6	0.30	0.30	
7	0.99	0.01	
8	1.00	0.01	
9	1.00	0.00	

	dst_host_same_src_port_rate	dst_host_srv_diff_host_rate	\
0	0.00	0.00	
1	0.00	0.00	
2	0.00	0.00	
3	0.00	0.00	
4	0.01	0.00	
5	0.01	0.00	
6	0.30	0.00	
7	0.00	0.00	
8	0.00	0.00	
9	0.01	0.01	

	dst_host_serror_rate	dst_host_srv_serror_rate	dst_host_rerror_rate	\
0	0.0	0.0	0.0	
1	0.0	0.0	0.0	
2	0.0	0.0	0.0	
3	0.0	0.0	0.0	
4	0.0	0.0	0.0	
5	0.0	0.0	0.0	
6	0.0	0.0	0.0	
7	0.0	0.0	0.0	
8	0.0	0.0	0.0	

9 0.0 0.0 0.0

	dst_host_srv_rerror_rate	attack_type
0	0.0	normal.
1	0.0	normal.
2	0.0	normal.
3	0.0	snmpgetattack.
4	0.0	snmpgetattack.
5	0.0	snmpgetattack.
6	0.0	normal.
7	0.0	normal.
8	0.0	snmpgetattack.
9	0.0	normal.

[10 rows x 42 columns]

In [4]: kdd_test_10percent.describe()

Out[4]:

	duration	src_bytes	dst_bytes	land	\
count	311029.000000	3.110290e+05	3.110290e+05	311029.000000	
mean	17.902736	1.731702e+03	7.479937e+02	0.000029	
std	407.644400	1.276567e+05	1.612018e+04	0.005379	
min	0.000000	0.000000e+00	0.000000e+00	0.000000	
25%	0.000000	1.050000e+02	0.000000e+00	0.000000	
50%	0.000000	5.200000e+02	0.000000e+00	0.000000	
75%	0.000000	1.032000e+03	0.000000e+00	0.000000	
max	57715.000000	6.282565e+07	5.203179e+06	1.000000	

	wrong_fragment	urgent	hot	num_failed_logins	\
count	311029.000000	311029.000000	311029.000000	311029.000000	
mean	0.000762	0.000051	0.014677	0.002363	
std	0.040367	0.009821	0.312068	0.049990	
min	0.000000	0.000000	0.000000	0.000000	
25%	0.000000	0.000000	0.000000	0.000000	
50%	0.000000	0.000000	0.000000	0.000000	
75%	0.000000	0.000000	0.000000	0.000000	
max	3.000000	3.000000	101.000000	4.000000	

	logged_in	num_compromised	...	\
count	311029.000000	311029.000000	...	
mean	0.172476	0.011243	...	
std	0.377794	1.958325	...	
min	0.000000	0.000000	...	
25%	0.000000	0.000000	...	
50%	0.000000	0.000000	...	
75%	0.000000	0.000000	...	
max	1.000000	796.000000	...	

	dst_host_count	dst_host_srv_count	dst_host_same_srv_rate \
count	311029.000000	311029.000000	311029.000000
mean	235.282681	199.193914	0.793494
std	60.913298	100.306470	0.387090
min	0.000000	0.000000	0.000000
25%	255.000000	244.000000	0.970000
50%	255.000000	255.000000	1.000000
75%	255.000000	255.000000	1.000000
max	255.000000	255.000000	1.000000

	dst_host_diff_srv_rate	dst_host_same_src_port_rate \
count	311029.000000	311029.000000
mean	0.024953	0.547919
std	0.096003	0.491963
min	0.000000	0.000000
25%	0.000000	0.000000
50%	0.000000	1.000000
75%	0.010000	1.000000
max	1.000000	1.000000

	dst_host_srv_diff_host_rate	dst_host_serror_rate \
count	311029.000000	311029.000000
mean	0.004566	0.058764
std	0.035773	0.231296
min	0.000000	0.000000
25%	0.000000	0.000000
50%	0.000000	0.000000
75%	0.000000	0.000000
max	1.000000	1.000000

	dst_host_srv_serror_rate	dst_host_rerror_rate \
count	311029.000000	311029.000000
mean	0.058791	0.142659
std	0.232997	0.344380
min	0.000000	0.000000
25%	0.000000	0.000000
50%	0.000000	0.000000
75%	0.000000	0.000000
max	1.000000	1.000000

	dst_host_srv_rerror_rate
count	311029.000000
mean	0.141693
std	0.346573
min	0.000000
25%	0.000000
50%	0.000000
75%	0.000000


```
max 1.000000
```

```
[8 rows x 38 columns]
```

1.1.2 Data Exploration

We now proceed with some basic data exploration on the training dataset. All modifications made on the training dataset (feature engineering) will be replicated on the test dataset (to ensure apples-to-apples comparison).

```
In [3]: import matplotlib.pyplot as plt
import numpy as np
import math
import seaborn as sns
```

```
%matplotlib inline
```

```
In [4]: from IPython.core.interactiveshell import InteractiveShell
InteractiveShell.ast_node_interactivity = "all"
```

```
In [7]: #Checking for null values in train and test datasets
```

```
kdd_train_10percent.isnull().any().any()
kdd_test_10percent.isnull().any().any()
```

```
#No nulls found
```

```
Out[7]: False
```

```
Out[7]: False
```

```
In [8]: #Comparing categorical variables across train and test datasets
```

```
len(kdd_train_10percent.protocol_type.unique())
sorted(kdd_train_10percent.protocol_type.unique())
```

```
Out[8]: 3
```

```
Out[8]: ['icmp', 'tcp', 'udp']
```

```
In [9]: len(kdd_test_10percent.protocol_type.unique())
sorted(kdd_test_10percent.protocol_type.unique())
```

```
Out[9]: 3
```

```
Out[9]: ['icmp', 'tcp', 'udp']
```

```
In [10]: len(kdd_train_10percent.service.unique())
sorted(kdd_train_10percent.service.unique())
```

Out[10]: 66

```
Out[10]: ['IRC',
          'X11',
          'Z39_50',
          'auth',
          'bgp',
          'courier',
          'csnet_ns',
          'ctf',
          'daytime',
          'discard',
          'domain',
          'domain_u',
          'echo',
          'eco_i',
          'ecr_i',
          'efs',
          'exec',
          'finger',
          'ftp',
          'ftp_data',
          'gopher',
          'hostnames',
          'http',
          'http_443',
          'imap4',
          'iso_tsap',
          'klogin',
          'kshell',
          'ldap',
          'link',
          'login',
          'mtp',
          'name',
          'netbios_dgm',
          'netbios_ns',
          'netbios_ssn',
          'netstat',
          'nnsp',
          'nntp',
          'ntp_u',
          'other',
          'pm_dump',
          'pop_2',
          'pop_3',
          'printer',
          'private',
```

```
'red_i',  
'remote_job',  
'rje',  
'shell',  
'smtp',  
'sql_net',  
'ssh',  
'sunrpc',  
'supdup',  
'sysstat',  
'telnet',  
'tftp_u',  
'tim_i',  
'time',  
'urh_i',  
'urp_i',  
'uucp',  
'uucp_path',  
'vmnet',  
'whois']
```

```
In [11]: len(kdd_test_10percent.service.unique())  
         sorted(kdd_test_10percent.service.unique())
```

```
Out[11]: 65
```

```
Out[11]: ['IRC',  
          'X11',  
          'Z39_50',  
          'auth',  
          'bgp',  
          'courier',  
          'csnet_ns',  
          'ctf',  
          'daytime',  
          'discard',  
          'domain',  
          'domain_u',  
          'echo',  
          'eco_i',  
          'ecr_i',  
          'efs',  
          'exec',  
          'finger',  
          'ftp',  
          'ftp_data',  
          'gopher',  
          'hostnames',
```

```
'http',
'http_443',
'icmp',
'imap4',
'iso_tsap',
'klogin',
'kshell',
'ldap',
'link',
'login',
'mtp',
'name',
'netbios_dgm',
'netbios_ns',
'netbios_ssn',
'netstat',
'nnspp',
'nntp',
'ntp_u',
'other',
'pm_dump',
'pop_2',
'pop_3',
'printer',
'private',
'remote_job',
'rje',
'shell',
'smtp',
'sql_net',
'ssh',
'sunrpc',
'supdup',
'sysstat',
'telnet',
'tftp_u',
'tim_i',
'time',
'urp_i',
'uucp',
'uucp_path',
'vmnet',
'whois']
```

```
In [12]: len(kdd_train_10percent.flag.unique())
         sorted(kdd_train_10percent.flag.unique())
```

```
Out[12]: 11
```

```
Out[12]: ['OTH', 'REJ', 'RSTO', 'RSTOSO', 'RSTR', 'SO', 'S1', 'S2', 'S3', 'SF', 'SH']
```

```
In [13]: len(kdd_test_10percent.flag.unique())
         sorted(kdd_test_10percent.flag.unique())
```

```
Out[13]: 11
```

```
Out[13]: ['OTH', 'REJ', 'RSTO', 'RSTOSO', 'RSTR', 'SO', 'S1', 'S2', 'S3', 'SF', 'SH']
```

We see that the "service" column has one less type in the test dataset than in the train dataset. We would not have been aware of this if we had not checked the test dataset. For simplicity, we remove the categorical variable columns for our analysis.

```
In [5]: #Removing the categorical variable columns
new_labels = ["duration", "src_bytes", "dst_bytes", "land", "wrong_fragment", "urgent", "hot",
              "logged_in", "num_compromised", "root_shell", "su_attempted", "num_root", "num_file_created",
              "is_host_login", "is_guest_login", "count", "srv_count", "serror_rate", "srv_serror_rate",
              "diff_srv_rate", "srv_diff_host_rate", "dst_host_count", "dst_host_srv_count", "dst_host_serror_rate",
              "dst_host_srv_diff_host_rate", "dst_host_serror_rate", "dst_host_srv_serror_rate", "dst_host_count"]
kdd_train_10per_mod = kdd_train_10percent[new_labels]
kdd_train_10per_mod.shape
kdd_train_10per_mod.describe()
```

```
Out[5]: (494021, 39)
```

```
Out[5]:
```

	duration	src_bytes	dst_bytes	land	\
count	494021.000000	4.940210e+05	4.940210e+05	494021.000000	
mean	47.979302	3.025610e+03	8.685324e+02	0.000045	
std	707.746472	9.882181e+05	3.304000e+04	0.006673	
min	0.000000	0.000000e+00	0.000000e+00	0.000000	
25%	0.000000	4.500000e+01	0.000000e+00	0.000000	
50%	0.000000	5.200000e+02	0.000000e+00	0.000000	
75%	0.000000	1.032000e+03	0.000000e+00	0.000000	
max	58329.000000	6.933756e+08	5.155468e+06	1.000000	

	wrong_fragment	urgent	hot	num_failed_logins	\
count	494021.000000	494021.000000	494021.000000	494021.000000	
mean	0.006433	0.000014	0.034519	0.000152	
std	0.134805	0.005510	0.782103	0.015520	
min	0.000000	0.000000	0.000000	0.000000	
25%	0.000000	0.000000	0.000000	0.000000	
50%	0.000000	0.000000	0.000000	0.000000	
75%	0.000000	0.000000	0.000000	0.000000	
max	3.000000	3.000000	30.000000	5.000000	

	logged_in	num_compromised	...	\
count	494021.000000	494021.000000	...	
mean	0.148247	0.010212	...	
std	0.355345	1.798326	...	

min	0.000000	0.000000	...
25%	0.000000	0.000000	...
50%	0.000000	0.000000	...
75%	0.000000	0.000000	...
max	1.000000	884.000000	...

	dst_host_count	dst_host_srv_count	dst_host_same_srv_rate \
count	494021.000000	494021.000000	494021.000000
mean	232.470778	188.665670	0.753780
std	64.745380	106.040437	0.410781
min	0.000000	0.000000	0.000000
25%	255.000000	46.000000	0.410000
50%	255.000000	255.000000	1.000000
75%	255.000000	255.000000	1.000000
max	255.000000	255.000000	1.000000

	dst_host_diff_srv_rate	dst_host_same_src_port_rate \
count	494021.000000	494021.000000
mean	0.030906	0.601935
std	0.109259	0.481309
min	0.000000	0.000000
25%	0.000000	0.000000
50%	0.000000	1.000000
75%	0.040000	1.000000
max	1.000000	1.000000

	dst_host_srv_diff_host_rate	dst_host_serror_rate \
count	494021.000000	494021.000000
mean	0.006684	0.176754
std	0.042133	0.380593
min	0.000000	0.000000
25%	0.000000	0.000000
50%	0.000000	0.000000
75%	0.000000	0.000000
max	1.000000	1.000000

	dst_host_srv_serror_rate	dst_host_rerror_rate \
count	494021.000000	494021.000000
mean	0.176443	0.058118
std	0.380919	0.230590
min	0.000000	0.000000
25%	0.000000	0.000000
50%	0.000000	0.000000
75%	0.000000	0.000000
max	1.000000	1.000000

	dst_host_srv_rerror_rate
count	494021.000000

```

mean          0.057412
std           0.230140
min           0.000000
25%           0.000000
50%           0.000000
75%           0.000000
max           1.000000

```

```
[8 rows x 38 columns]
```

```
In [6]: #Do the same for test data to keep the columns consistent
```

```

kdd_test_10per_mod = kdd_test_10percent[new_labels]
kdd_test_10per_mod.shape
kdd_test_10per_mod.describe()

```

```
Out[6]: (311029, 39)
```

```

Out[6]:
      duration  src_bytes  dst_bytes  land \
count  311029.000000  3.110290e+05  3.110290e+05  311029.000000
mean      17.902736  1.731702e+03  7.479937e+02    0.000029
std      407.644400  1.276567e+05  1.612018e+04    0.005379
min        0.000000  0.000000e+00  0.000000e+00    0.000000
25%        0.000000  1.050000e+02  0.000000e+00    0.000000
50%        0.000000  5.200000e+02  0.000000e+00    0.000000
75%        0.000000  1.032000e+03  0.000000e+00    0.000000
max      57715.000000  6.282565e+07  5.203179e+06    1.000000

      wrong_fragment  urgent  hot  num_failed_logins \
count  311029.000000  311029.000000  311029.000000  311029.000000
mean      0.000762    0.000051    0.014677    0.002363
std      0.040367    0.009821    0.312068    0.049990
min        0.000000    0.000000    0.000000    0.000000
25%        0.000000    0.000000    0.000000    0.000000
50%        0.000000    0.000000    0.000000    0.000000
75%        0.000000    0.000000    0.000000    0.000000
max         3.000000    3.000000   101.000000    4.000000

      logged_in  num_compromised  ... \
count  311029.000000  311029.000000  ...
mean      0.172476    0.011243  ...
std      0.377794    1.958325  ...
min        0.000000    0.000000  ...
25%        0.000000    0.000000  ...
50%        0.000000    0.000000  ...
75%        0.000000    0.000000  ...
max         1.000000   796.000000  ...

```

	dst_host_count	dst_host_srv_count	dst_host_same_srv_rate \
count	311029.000000	311029.000000	311029.000000
mean	235.282681	199.193914	0.793494
std	60.913298	100.306470	0.387090
min	0.000000	0.000000	0.000000
25%	255.000000	244.000000	0.970000
50%	255.000000	255.000000	1.000000
75%	255.000000	255.000000	1.000000
max	255.000000	255.000000	1.000000

	dst_host_diff_srv_rate	dst_host_same_src_port_rate \
count	311029.000000	311029.000000
mean	0.024953	0.547919
std	0.096003	0.491963
min	0.000000	0.000000
25%	0.000000	0.000000
50%	0.000000	1.000000
75%	0.010000	1.000000
max	1.000000	1.000000

	dst_host_srv_diff_host_rate	dst_host_serror_rate \
count	311029.000000	311029.000000
mean	0.004566	0.058764
std	0.035773	0.231296
min	0.000000	0.000000
25%	0.000000	0.000000
50%	0.000000	0.000000
75%	0.000000	0.000000
max	1.000000	1.000000

	dst_host_srv_serror_rate	dst_host_rerror_rate \
count	311029.000000	311029.000000
mean	0.058791	0.142659
std	0.232997	0.344380
min	0.000000	0.000000
25%	0.000000	0.000000
50%	0.000000	0.000000
75%	0.000000	0.000000
max	1.000000	1.000000

	dst_host_srv_rerror_rate
count	311029.000000
mean	0.141693
std	0.346573
min	0.000000
25%	0.000000
50%	0.000000
75%	0.000000


```
max 1.000000
```

```
[8 rows x 38 columns]
```

```
In [7]: #Exploring the target variable
```

```
#According to the data documentation the attack types can be divided into "normal" and 4
```

```
attack_labs_all = kdd_train_10per_mod['attack_type'].copy()
len(attack_labs_all.unique())
```

```
dos_list = ["back.", "land.", "neptune.", "pod.", "smurf.", "teardrop."]
utr_list = ["buffer_overflow.", "loadmodule.", "perl.", "rootkit."]
rtl_list = ["ftp_write.", "guess_passwd.", "imap.", "multihop.", "phf.", "spy.", "warezclient."]
probe_list = ["ipsweep.", "nmap.", "portsweep.", "satan."]
```

```
attack_labs_all[attack_labs_all.isin(dos_list)] = "DOS."
attack_labs_all[attack_labs_all.isin(utr_list)] = "U2R"
attack_labs_all[attack_labs_all.isin(rtl_list)] = "R2L."
attack_labs_all[attack_labs_all.isin(probe_list)] = "probe."
```

```
df = attack_labs_all.value_counts()
print(df)
plt.figure()
df.plot(kind='bar')
plt.xlabel("Attack Types in Training Data")
plt.ylabel("Value Counts")
```

```
#We see that majority of the attacks are DOS attacks (98.67 percent)
#So for simplicity we will just classify the labels as "normal." and "bad." attacks
```

```
#Checking in the test dataset as well (has extra attack types too)
```

```
attack_labs_tall = kdd_test_10per_mod['attack_type'].copy()
len(attack_labs_tall.unique())
```

```
dos_list = ["back.", "land.", "neptune.", "pod.", "smurf.", "teardrop.", "apache2.", "mailbomb."]
utr_list = ["buffer_overflow.", "loadmodule.", "perl.", "rootkit.", "ps.", "sqlattack.", "xterm."]
rtl_list = ["ftp_write.", "guess_passwd.", "imap.", "multihop.", "phf.", "spy.", "warezclient.",
            "sendmail.", "snmpgetattack.", "snmpguess.", "worm.", "xlock.", "xsnoop."]
probe_list = ["ipsweep.", "nmap.", "portsweep.", "satan.", "mscan.", "saint."]
```

```
attack_labs_tall[attack_labs_tall.isin(dos_list)] = "DOS."
attack_labs_tall[attack_labs_tall.isin(utr_list)] = "U2R"
attack_labs_tall[attack_labs_tall.isin(rtl_list)] = "R2L."
attack_labs_tall[attack_labs_tall.isin(probe_list)] = "probe."
```

```
df = attack_labs_tall.value_counts()
print(df)
plt.figure()
```

```
df.plot(kind='bar')
plt.xlabel("Attack Types in Test Data")
plt.ylabel("Value Counts")
```

#Here too we see extra attack types but maximum of the labels (out of all attack types)

Out[7]: 23

```
DOS.      391458
normal.   97278
probe.    4107
R2L.      1126
U2R       52
Name: attack_type, dtype: int64
```

Out[7]: <matplotlib.figure.Figure at 0x221cd5c54a8>

Out[7]: <matplotlib.axes._subplots.AxesSubplot at 0x221cd5c5c50>

Out[7]: Text(0.5,0,'Attack Types in Training Data')

Out[7]: Text(0,0.5,'Value Counts')

Out[7]: 38

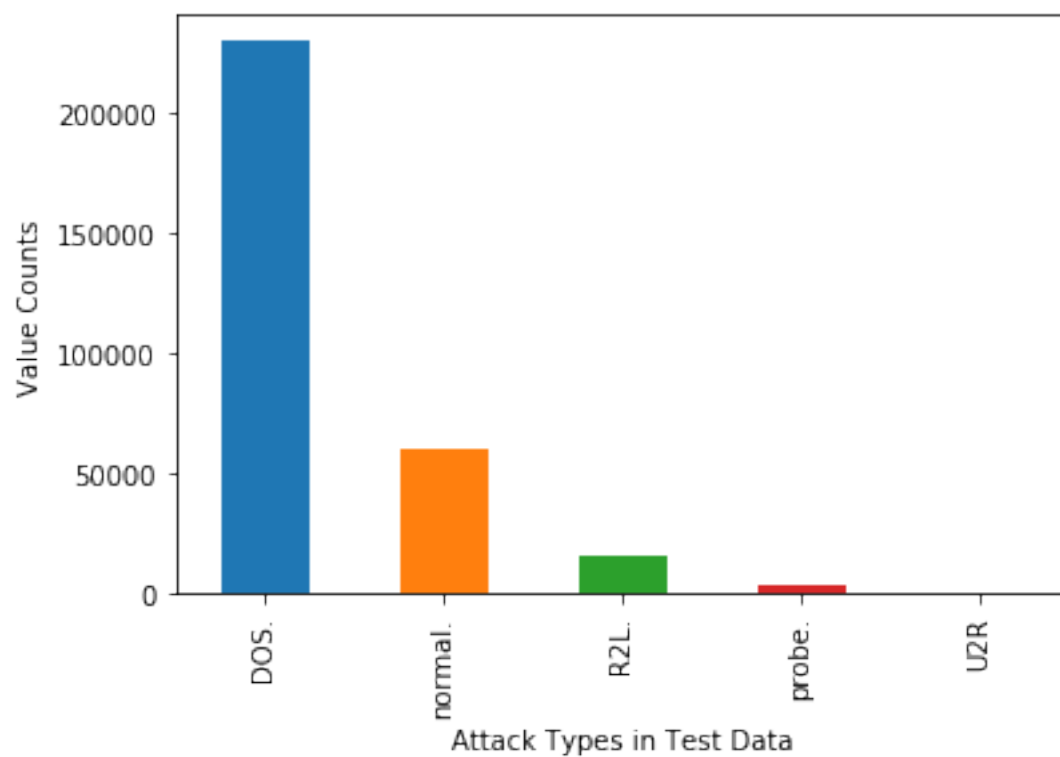
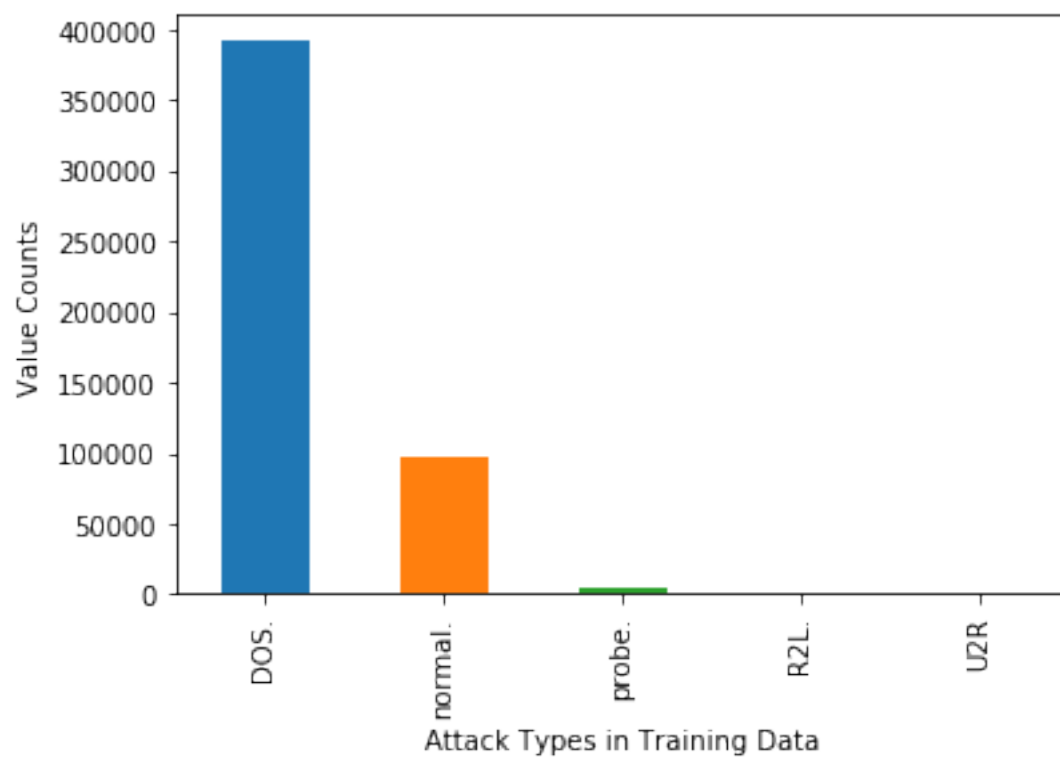
```
DOS.      229853
normal.   60593
R2L.      16189
probe.    4166
U2R       228
Name: attack_type, dtype: int64
```

Out[7]: <matplotlib.figure.Figure at 0x221cdf1f630>

Out[7]: <matplotlib.axes._subplots.AxesSubplot at 0x221cdebed68>

Out[7]: Text(0.5,0,'Attack Types in Test Data')

Out[7]: Text(0,0.5,'Value Counts')



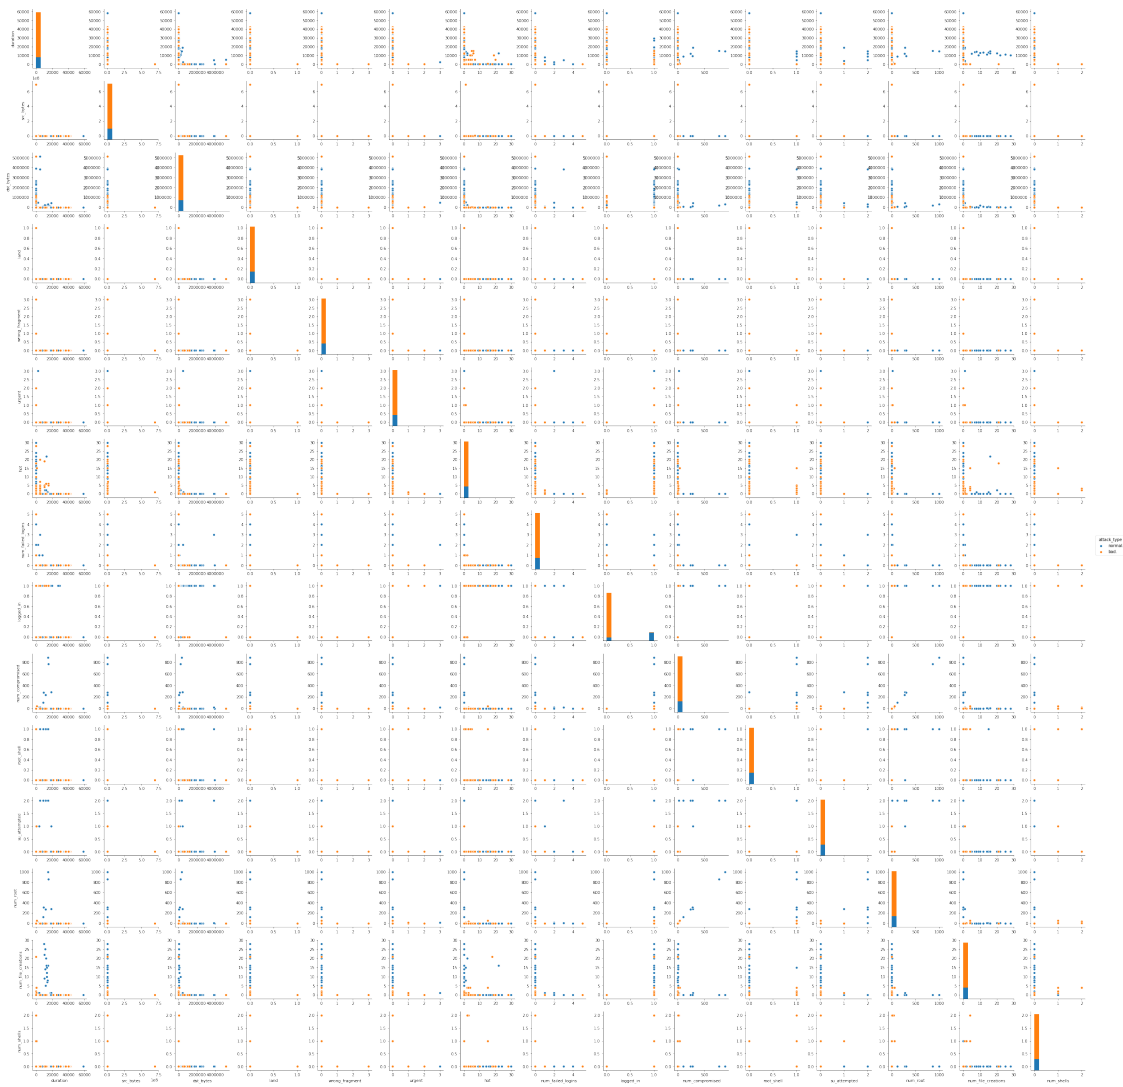
```
In [13]: #change the labels to only show 'normal' and 'bad' as the attack types
attack_labs = kdd_train_10per_mod['attack_type'].copy()
attack_labs[attack_labs!='normal.'] = 'bad.'
attack_labs.value_counts()
```

```
Out[13]: bad.          396743
normal.         97278
Name: attack_type, dtype: int64
```

```
In [21]: #Creating pairplots for some continuous variables (first 15) in the training dataset
train_data_for_pairplot = pd.concat([kdd_train_10per_mod.iloc[:,0:15],attack_labs], axis=1)

#Pairplot
var_plot = sns.pairplot(data = train_data_for_pairplot, hue = "attack_type", kind='scatter')
plt.show()
#var_plot.savefig("pairplot.png")
```

```
Out[21]: bad.          396743
normal.         97278
Name: attack_type, dtype: int64
```



Even though we have plotted the distributions and correlations for the first few variables, we can make two important observations: the data seems skewed and there do not seem to be a lot of strong linear correlations visible between these variables. Since we will be using machine learning techniques like K-Nearest Neighbors (involving distance metrics) along with other algorithms and will also be verifying results using clustering (unsupervised check), we should scale the continuous variables.

1.1.3 Data Preprocessing : Scaling

In this step, we scale the continuous variables in our training dataset as part of the data preprocessing. We will later use the same transformation for the test data.

In [8]: *#Creating train dataset without the target column for scaling*

```
kdd_train_10per_var = kdd_train_10per_mod.iloc[:,0:38]
```

```

kdd_train_10per_var.shape
kdd_train_10per_var.describe()
kdd_train_10per_var.head(10)
type(kdd_train_10per_var)

#Scaling only for continuous variables -- so we remove the discrete variables
#Creating a list of the discrete valued columns (mentioned in the documentation for the
cont_labels = ["duration", "src_bytes", "dst_bytes", "wrong_fragment", "urgent", "hot", "num_
               "num_compromised", "num_root", "num_file_creations", "num_shells", "num_access_files", "n
               "count", "srv_count", "error_rate", "srv_error_rate", "error_rate", "srv_error_rate",
               "diff_srv_rate", "srv_diff_host_rate", "dst_host_count", "dst_host_srv_count", "dst_host
               "dst_host_srv_diff_host_rate", "dst_host_error_rate", "dst_host_srv_error_rate", "dst
other_labels = ["land", "logged_in", "root_shell", "su_attempted", "is_host_login", "is_guest
kdd_train_10per_forscale = kdd_train_10per_mod[cont_labels]
kdd_train_10per_forscale.shape
kdd_train_10per_forscale.describe()

kdd_train_10per_others = kdd_train_10per_mod[other_labels]
kdd_train_10per_others.shape
kdd_train_10per_others.describe()

```

Out[8]: (494021, 38)

```

Out[8]:

```

	duration	src_bytes	dst_bytes	land	\
count	494021.000000	4.940210e+05	4.940210e+05	494021.000000	
mean	47.979302	3.025610e+03	8.685324e+02	0.000045	
std	707.746472	9.882181e+05	3.304000e+04	0.006673	
min	0.000000	0.000000e+00	0.000000e+00	0.000000	
25%	0.000000	4.500000e+01	0.000000e+00	0.000000	
50%	0.000000	5.200000e+02	0.000000e+00	0.000000	
75%	0.000000	1.032000e+03	0.000000e+00	0.000000	
max	58329.000000	6.933756e+08	5.155468e+06	1.000000	

	wrong_fragment	urgent	hot	num_failed_logins	\
count	494021.000000	494021.000000	494021.000000	494021.000000	
mean	0.006433	0.000014	0.034519	0.000152	
std	0.134805	0.005510	0.782103	0.015520	
min	0.000000	0.000000	0.000000	0.000000	
25%	0.000000	0.000000	0.000000	0.000000	
50%	0.000000	0.000000	0.000000	0.000000	
75%	0.000000	0.000000	0.000000	0.000000	
max	3.000000	3.000000	30.000000	5.000000	

	logged_in	num_compromised	...	\
count	494021.000000	494021.000000	...	
mean	0.148247	0.010212	...	
std	0.355345	1.798326	...	
min	0.000000	0.000000	...	

25%	0.000000	0.000000	...
50%	0.000000	0.000000	...
75%	0.000000	0.000000	...
max	1.000000	884.000000	...

	dst_host_count	dst_host_srv_count	dst_host_same_srv_rate \
count	494021.000000	494021.000000	494021.000000
mean	232.470778	188.665670	0.753780
std	64.745380	106.040437	0.410781
min	0.000000	0.000000	0.000000
25%	255.000000	46.000000	0.410000
50%	255.000000	255.000000	1.000000
75%	255.000000	255.000000	1.000000
max	255.000000	255.000000	1.000000

	dst_host_diff_srv_rate	dst_host_same_src_port_rate \
count	494021.000000	494021.000000
mean	0.030906	0.601935
std	0.109259	0.481309
min	0.000000	0.000000
25%	0.000000	0.000000
50%	0.000000	1.000000
75%	0.040000	1.000000
max	1.000000	1.000000

	dst_host_srv_diff_host_rate	dst_host_serror_rate \
count	494021.000000	494021.000000
mean	0.006684	0.176754
std	0.042133	0.380593
min	0.000000	0.000000
25%	0.000000	0.000000
50%	0.000000	0.000000
75%	0.000000	0.000000
max	1.000000	1.000000

	dst_host_srv_serror_rate	dst_host_rerror_rate \
count	494021.000000	494021.000000
mean	0.176443	0.058118
std	0.380919	0.230590
min	0.000000	0.000000
25%	0.000000	0.000000
50%	0.000000	0.000000
75%	0.000000	0.000000
max	1.000000	1.000000

	dst_host_srv_rerror_rate
count	494021.000000
mean	0.057412

```

std          0.230140
min          0.000000
25%          0.000000
50%          0.000000
75%          0.000000
max          1.000000

```

[8 rows x 38 columns]

```

Out[8]:      duration  src_bytes  dst_bytes  land  wrong_fragment  urgent  hot  \
0           0         181         5450    0           0           0    0
1           0         239          486    0           0           0    0
2           0         235         1337    0           0           0    0
3           0         219         1337    0           0           0    0
4           0         217         2032    0           0           0    0
5           0         217         2032    0           0           0    0
6           0         212         1940    0           0           0    0
7           0         159         4087    0           0           0    0
8           0         210          151    0           0           0    0
9           0         212          786    0           0           0    1

```

```

      num_failed_logins  logged_in  num_compromised  ...  \
0                   0           1           0       ...
1                   0           1           0       ...
2                   0           1           0       ...
3                   0           1           0       ...
4                   0           1           0       ...
5                   0           1           0       ...
6                   0           1           0       ...
7                   0           1           0       ...
8                   0           1           0       ...
9                   0           1           0       ...

```

```

      dst_host_count  dst_host_srv_count  dst_host_same_srv_rate  \
0                   9                  9                  1.0
1                  19                 19                  1.0
2                  29                 29                  1.0
3                  39                 39                  1.0
4                  49                 49                  1.0
5                  59                 59                  1.0
6                   1                 69                  1.0
7                  11                 79                  1.0
8                   8                 89                  1.0
9                   8                 99                  1.0

```

```

      dst_host_diff_srv_rate  dst_host_same_src_port_rate  \
0                   0.0                   0.11
1                   0.0                   0.05

```


2	0.0	0.03
3	0.0	0.03
4	0.0	0.02
5	0.0	0.02
6	0.0	1.00
7	0.0	0.09
8	0.0	0.12
9	0.0	0.12

	dst_host_srv_diff_host_rate	dst_host_serror_rate \
0	0.00	0.0
1	0.00	0.0
2	0.00	0.0
3	0.00	0.0
4	0.00	0.0
5	0.00	0.0
6	0.04	0.0
7	0.04	0.0
8	0.04	0.0
9	0.05	0.0

	dst_host_srv_serror_rate	dst_host_rerror_rate	dst_host_srv_rerror_rate
0	0.0	0.0	0.0
1	0.0	0.0	0.0
2	0.0	0.0	0.0
3	0.0	0.0	0.0
4	0.0	0.0	0.0
5	0.0	0.0	0.0
6	0.0	0.0	0.0
7	0.0	0.0	0.0
8	0.0	0.0	0.0
9	0.0	0.0	0.0

[10 rows x 38 columns]

Out[8]: pandas.core.frame.DataFrame

Out[8]: (494021, 32)

Out[8]:

	duration	src_bytes	dst_bytes	wrong_fragment \
count	494021.000000	4.940210e+05	4.940210e+05	494021.000000
mean	47.979302	3.025610e+03	8.685324e+02	0.006433
std	707.746472	9.882181e+05	3.304000e+04	0.134805
min	0.000000	0.000000e+00	0.000000e+00	0.000000
25%	0.000000	4.500000e+01	0.000000e+00	0.000000
50%	0.000000	5.200000e+02	0.000000e+00	0.000000
75%	0.000000	1.032000e+03	0.000000e+00	0.000000
max	58329.000000	6.933756e+08	5.155468e+06	3.000000

	urgent	hot	num_failed_logins	num_compromised	\
count	494021.000000	494021.000000	494021.000000	494021.000000	
mean	0.000014	0.034519	0.000152	0.010212	
std	0.005510	0.782103	0.015520	1.798326	
min	0.000000	0.000000	0.000000	0.000000	
25%	0.000000	0.000000	0.000000	0.000000	
50%	0.000000	0.000000	0.000000	0.000000	
75%	0.000000	0.000000	0.000000	0.000000	
max	3.000000	30.000000	5.000000	884.000000	

	num_root	num_file_creations	...	\
count	494021.000000	494021.000000	...	
mean	0.011352	0.001083	...	
std	2.012718	0.096416	...	
min	0.000000	0.000000	...	
25%	0.000000	0.000000	...	
50%	0.000000	0.000000	...	
75%	0.000000	0.000000	...	
max	993.000000	28.000000	...	

	dst_host_count	dst_host_srv_count	dst_host_same_srv_rate	\
count	494021.000000	494021.000000	494021.000000	
mean	232.470778	188.665670	0.753780	
std	64.745380	106.040437	0.410781	
min	0.000000	0.000000	0.000000	
25%	255.000000	46.000000	0.410000	
50%	255.000000	255.000000	1.000000	
75%	255.000000	255.000000	1.000000	
max	255.000000	255.000000	1.000000	

	dst_host_diff_srv_rate	dst_host_same_src_port_rate	\
count	494021.000000	494021.000000	
mean	0.030906	0.601935	
std	0.109259	0.481309	
min	0.000000	0.000000	
25%	0.000000	0.000000	
50%	0.000000	1.000000	
75%	0.040000	1.000000	
max	1.000000	1.000000	

	dst_host_srv_diff_host_rate	dst_host_serror_rate	\
count	494021.000000	494021.000000	
mean	0.006684	0.176754	
std	0.042133	0.380593	
min	0.000000	0.000000	
25%	0.000000	0.000000	
50%	0.000000	0.000000	

75%	0.000000	0.000000
max	1.000000	1.000000

	dst_host_srv_serror_rate	dst_host_rerror_rate \
count	494021.000000	494021.000000
mean	0.176443	0.058118
std	0.380919	0.230590
min	0.000000	0.000000
25%	0.000000	0.000000
50%	0.000000	0.000000
75%	0.000000	0.000000
max	1.000000	1.000000

	dst_host_srv_rerror_rate
count	494021.000000
mean	0.057412
std	0.230140
min	0.000000
25%	0.000000
50%	0.000000
75%	0.000000
max	1.000000

[8 rows x 32 columns]

Out[8]: (494021, 6)

	land	logged_in	root_shell	su_attempted \
count	494021.000000	494021.000000	494021.000000	494021.000000
mean	0.000045	0.148247	0.000111	0.000036
std	0.006673	0.355345	0.010551	0.007793
min	0.000000	0.000000	0.000000	0.000000
25%	0.000000	0.000000	0.000000	0.000000
50%	0.000000	0.000000	0.000000	0.000000
75%	0.000000	0.000000	0.000000	0.000000
max	1.000000	1.000000	1.000000	2.000000

	is_host_login	is_guest_login
count	494021.0	494021.000000
mean	0.0	0.001387
std	0.0	0.037211
min	0.0	0.000000
25%	0.0	0.000000
50%	0.0	0.000000
75%	0.0	0.000000
max	0.0	1.000000

In [9]: *#Now we scale the training dataset*

```

from sklearn import preprocessing

#Keeping the scaling parameters from the training dataset to be later applied to the test
train_scaler = preprocessing.StandardScaler().fit(kdd_train_10per_for)
train_scaler

train_scaler.mean_
train_scaler.scale_

#scaling the train dataset
kdd_train_10per_scaled1 = pd.DataFrame(data = train_scaler.transform(kdd_train_10per_for))
kdd_train_10per_scaled1.shape
kdd_train_10per_scaled1.describe()
kdd_train_10per_scaled1.head(10)

#adding the binary valued variable columns
kdd_train_10per_scaled = pd.concat([kdd_train_10per_scaled1, kdd_train_10per_others], axis=1)
kdd_train_10per_scaled.shape
kdd_train_10per_scaled.describe()
kdd_train_10per_scaled.head(10)

```

```
Out[9]: StandardScaler(copy=True, with_mean=True, with_std=True)
```

```
Out[9]: array([[ 4.79793025e+01,  3.02561030e+03,  8.68532425e+02,
  6.43292492e-03,  1.41694381e-05,  3.45187755e-02,
  1.51815409e-04,  1.02121165e-02,  1.13517442e-02,
  1.08294992e-03,  1.09307094e-04,  1.00805431e-03,
  0.00000000e+00,  3.32285690e+02,  2.92906557e+02,
  1.76686659e-01,  1.76608808e-01,  5.74334087e-02,
  5.77189431e-02,  7.91547343e-01,  2.09823874e-02,
  2.89968038e-02,  2.32470778e+02,  1.88665670e+02,
  7.53779698e-01,  3.09057307e-02,  6.01934756e-01,
  6.68350131e-03,  1.76753964e-01,  1.76442621e-01,
  5.81176104e-02,  5.74116687e-02])
```

```
Out[9]: array([[ 7.07745756e+02,  9.88217101e+05,  3.30399678e+04,
  1.34805112e-01,  5.51025235e-03,  7.82101790e-01,
  1.55195812e-02,  1.79832444e+00,  2.01271629e+00,
  9.64157811e-02,  1.10199990e-02,  3.64816531e-02,
  1.00000000e+00,  2.13147196e+02,  2.46322568e+02,
  3.80716571e-01,  3.81016199e-01,  2.31623239e-01,
  2.32146746e-01,  3.88189100e-01,  8.22054097e-02,
  1.42397323e-01,  6.47453148e+01,  1.06040330e+02,
  4.10780562e-01,  1.09259002e-01,  4.81308764e-01,
  4.21328318e-02,  3.80592712e-01,  3.80919063e-01,
  2.30589274e-01,  2.30140091e-01])
```

```
Out[9]: (494021, 32)
```

```

Out[9]:
      duration      src_bytes      dst_bytes  wrong_fragment      urgent \
count  4.940210e+05  4.940210e+05  4.940210e+05    4.940210e+05  4.940210e+05
mean  -3.625574e-14 -1.320064e-14 -5.498890e-14   -6.073146e-15 -1.529804e-14
std    1.000001e+00  1.000001e+00  1.000001e+00    1.000001e+00  1.000001e+00
min    -6.779172e-02 -3.061686e-03 -2.628733e-02   -4.772019e-02 -2.571468e-03
25%    -6.779172e-02 -3.016149e-03 -2.628733e-02   -4.772019e-02 -2.571468e-03
50%    -6.779172e-02 -2.535486e-03 -2.628733e-02   -4.772019e-02 -2.571468e-03
75%    -6.779172e-02 -2.017381e-03 -2.628733e-02   -4.772019e-02 -2.571468e-03
max     8.234740e+01  7.016400e+02  1.560110e+02    2.220663e+01  5.444371e+02

      hot  num_failed_logins  num_compromised      num_root \
count  4.940210e+05      4.940210e+05      4.940210e+05  4.940210e+05
mean  -6.957988e-14      -8.038565e-14     -1.113045e-14 -1.364308e-14
std    1.000001e+00      1.000001e+00      1.000001e+00  1.000001e+00
min    -4.413591e-02      -9.782185e-03     -5.678684e-03 -5.640012e-03
25%    -4.413591e-02      -9.782185e-03     -5.678684e-03 -5.640012e-03
50%    -4.413591e-02      -9.782185e-03     -5.678684e-03 -5.640012e-03
75%    -4.413591e-02      -9.782185e-03     -5.678684e-03 -5.640012e-03
max     3.831404e+01      3.221639e+02      4.915630e+02  4.933575e+02

      num_file_creations      ...      dst_host_count \
count      4.940210e+05      ...      4.940210e+05
mean      1.653658e-14      ...      -7.353653e-13
std        1.000001e+00      ...      1.000001e+00
min       -1.123208e-02      ...      -3.590542e+00
25%       -1.123208e-02      ...      3.479668e-01
50%       -1.123208e-02      ...      3.479668e-01
75%       -1.123208e-02      ...      3.479668e-01
max        2.903977e+02      ...      3.479668e-01

      dst_host_srv_count  dst_host_same_srv_rate  dst_host_diff_srv_rate \
count      4.940210e+05      4.940210e+05      4.940210e+05
mean      1.953080e-13      3.799918e-13      -1.392772e-13
std        1.000001e+00      1.000001e+00      1.000001e+00
min       -1.779188e+00      -1.834994e+00      -2.828667e-01
25%       -1.345391e+00      -8.368938e-01      -2.828667e-01
50%        6.255576e-01      5.993962e-01      -2.828667e-01
75%        6.255576e-01      5.993962e-01      8.323588e-02
max        6.255576e-01      5.993962e-01      8.869697e+00

      dst_host_same_src_port_rate  dst_host_srv_diff_host_rate \
count      4.940210e+05      4.940210e+05
mean      7.581150e-13      2.700207e-13
std        1.000001e+00      1.000001e+00
min       -1.250621e+00      -1.586293e-01
25%       -1.250621e+00      -1.586293e-01
50%        8.270476e-01      -1.586293e-01
75%        8.270476e-01      -1.586293e-01

```

max	8.270476e-01	2.357583e+01
-----	--------------	--------------

	dst_host_serror_rate	dst_host_srv_serror_rate	dst_host_rerror_rate \
count	4.940210e+05	4.940210e+05	4.940210e+05
mean	2.263813e-13	3.603748e-13	6.990594e-13
std	1.000001e+00	1.000001e+00	1.000001e+00
min	-4.644176e-01	-4.632024e-01	-2.520395e-01
25%	-4.644176e-01	-4.632024e-01	-2.520395e-01
50%	-4.644176e-01	-4.632024e-01	-2.520395e-01
75%	-4.644176e-01	-4.632024e-01	-2.520395e-01
max	2.163063e+00	2.162027e+00	4.084676e+00

	dst_host_srv_rerror_rate
count	4.940210e+05
mean	-9.654619e-14
std	1.000001e+00
min	-2.494640e-01
25%	-2.494640e-01
50%	-2.494640e-01
75%	-2.494640e-01
max	4.095715e+00

[8 rows x 32 columns]

Out[9]:

	duration	src_bytes	dst_bytes	wrong_fragment	urgent	hot \
0	-0.067792	-0.002879	0.138664	-0.04772	-0.002571	-0.044136
1	-0.067792	-0.002820	-0.011578	-0.04772	-0.002571	-0.044136
2	-0.067792	-0.002824	0.014179	-0.04772	-0.002571	-0.044136
3	-0.067792	-0.002840	0.014179	-0.04772	-0.002571	-0.044136
4	-0.067792	-0.002842	0.035214	-0.04772	-0.002571	-0.044136
5	-0.067792	-0.002842	0.035214	-0.04772	-0.002571	-0.044136
6	-0.067792	-0.002847	0.032429	-0.04772	-0.002571	-0.044136
7	-0.067792	-0.002901	0.097411	-0.04772	-0.002571	-0.044136
8	-0.067792	-0.002849	-0.021717	-0.04772	-0.002571	-0.044136
9	-0.067792	-0.002847	-0.002498	-0.04772	-0.002571	1.234470

	num_failed_logins	num_compromised	num_root	num_file_creations \
0	-0.009782	-0.005679	-0.00564	-0.011232
1	-0.009782	-0.005679	-0.00564	-0.011232
2	-0.009782	-0.005679	-0.00564	-0.011232
3	-0.009782	-0.005679	-0.00564	-0.011232
4	-0.009782	-0.005679	-0.00564	-0.011232
5	-0.009782	-0.005679	-0.00564	-0.011232
6	-0.009782	-0.005679	-0.00564	-0.011232
7	-0.009782	-0.005679	-0.00564	-0.011232
8	-0.009782	-0.005679	-0.00564	-0.011232
9	-0.009782	-0.005679	-0.00564	-0.011232

	...	dst_host_count	dst_host_srv_count	\
0	...	-3.451536	-1.694315	
1	...	-3.297085	-1.600011	
2	...	-3.142633	-1.505707	
3	...	-2.988182	-1.411403	
4	...	-2.833731	-1.317100	
5	...	-2.679279	-1.222796	
6	...	-3.575097	-1.128492	
7	...	-3.420646	-1.034188	
8	...	-3.466981	-0.939885	
9	...	-3.466981	-0.845581	

	dst_host_same_srv_rate	dst_host_diff_srv_rate	\
0	0.599396	-0.282867	
1	0.599396	-0.282867	
2	0.599396	-0.282867	
3	0.599396	-0.282867	
4	0.599396	-0.282867	
5	0.599396	-0.282867	
6	0.599396	-0.282867	
7	0.599396	-0.282867	
8	0.599396	-0.282867	
9	0.599396	-0.282867	

	dst_host_same_src_port_rate	dst_host_srv_diff_host_rate	\
0	-1.022077	-0.158629	
1	-1.146737	-0.158629	
2	-1.188291	-0.158629	
3	-1.188291	-0.158629	
4	-1.209067	-0.158629	
5	-1.209067	-0.158629	
6	0.827048	0.790749	
7	-1.063631	0.790749	
8	-1.001301	0.790749	
9	-1.001301	1.028094	

	dst_host_serror_rate	dst_host_srv_serror_rate	dst_host_rerror_rate	\
0	-0.464418	-0.463202	-0.25204	
1	-0.464418	-0.463202	-0.25204	
2	-0.464418	-0.463202	-0.25204	
3	-0.464418	-0.463202	-0.25204	
4	-0.464418	-0.463202	-0.25204	
5	-0.464418	-0.463202	-0.25204	
6	-0.464418	-0.463202	-0.25204	
7	-0.464418	-0.463202	-0.25204	
8	-0.464418	-0.463202	-0.25204	
9	-0.464418	-0.463202	-0.25204	

```

dst_host_srv_rerror_rate
0          -0.249464
1          -0.249464
2          -0.249464
3          -0.249464
4          -0.249464
5          -0.249464
6          -0.249464
7          -0.249464
8          -0.249464
9          -0.249464

```

[10 rows x 32 columns]

Out[9]: (494021, 38)

Out[9]:

	duration	src_bytes	dst_bytes	wrong_fragment	urgent \
count	4.940210e+05	4.940210e+05	4.940210e+05	4.940210e+05	4.940210e+05
mean	-3.625574e-14	-1.320064e-14	-5.498890e-14	-6.073146e-15	-1.529804e-14
std	1.000001e+00	1.000001e+00	1.000001e+00	1.000001e+00	1.000001e+00
min	-6.779172e-02	-3.061686e-03	-2.628733e-02	-4.772019e-02	-2.571468e-03
25%	-6.779172e-02	-3.016149e-03	-2.628733e-02	-4.772019e-02	-2.571468e-03
50%	-6.779172e-02	-2.535486e-03	-2.628733e-02	-4.772019e-02	-2.571468e-03
75%	-6.779172e-02	-2.017381e-03	-2.628733e-02	-4.772019e-02	-2.571468e-03
max	8.234740e+01	7.016400e+02	1.560110e+02	2.220663e+01	5.444371e+02

	hot	num_failed_logins	num_compromised	num_root \
count	4.940210e+05	4.940210e+05	4.940210e+05	4.940210e+05
mean	-6.957988e-14	-8.038565e-14	-1.113045e-14	-1.364308e-14
std	1.000001e+00	1.000001e+00	1.000001e+00	1.000001e+00
min	-4.413591e-02	-9.782185e-03	-5.678684e-03	-5.640012e-03
25%	-4.413591e-02	-9.782185e-03	-5.678684e-03	-5.640012e-03
50%	-4.413591e-02	-9.782185e-03	-5.678684e-03	-5.640012e-03
75%	-4.413591e-02	-9.782185e-03	-5.678684e-03	-5.640012e-03
max	3.831404e+01	3.221639e+02	4.915630e+02	4.933575e+02

	num_file_creations	...	dst_host_serror_rate \
count	4.940210e+05	...	4.940210e+05
mean	1.653658e-14	...	2.263813e-13
std	1.000001e+00	...	1.000001e+00
min	-1.123208e-02	...	-4.644176e-01
25%	-1.123208e-02	...	-4.644176e-01
50%	-1.123208e-02	...	-4.644176e-01
75%	-1.123208e-02	...	-4.644176e-01
max	2.903977e+02	...	2.163063e+00

	dst_host_srv_serror_rate	dst_host_rerror_rate \
count	4.940210e+05	4.940210e+05

mean	3.603748e-13	6.990594e-13
std	1.000001e+00	1.000001e+00
min	-4.632024e-01	-2.520395e-01
25%	-4.632024e-01	-2.520395e-01
50%	-4.632024e-01	-2.520395e-01
75%	-4.632024e-01	-2.520395e-01
max	2.162027e+00	4.084676e+00

	dst_host_srv_rerror_rate	land	logged_in	root_shell	\
count	4.940210e+05	494021.000000	494021.000000	494021.000000	
mean	-9.654619e-14	0.000045	0.148247	0.000111	
std	1.000001e+00	0.006673	0.355345	0.010551	
min	-2.494640e-01	0.000000	0.000000	0.000000	
25%	-2.494640e-01	0.000000	0.000000	0.000000	
50%	-2.494640e-01	0.000000	0.000000	0.000000	
75%	-2.494640e-01	0.000000	0.000000	0.000000	
max	4.095715e+00	1.000000	1.000000	1.000000	

	su_attempted	is_host_login	is_guest_login
count	494021.000000	494021.0	494021.000000
mean	0.000036	0.0	0.001387
std	0.007793	0.0	0.037211
min	0.000000	0.0	0.000000
25%	0.000000	0.0	0.000000
50%	0.000000	0.0	0.000000
75%	0.000000	0.0	0.000000
max	2.000000	0.0	1.000000

[8 rows x 38 columns]

Out[9]:

	duration	src_bytes	dst_bytes	wrong_fragment	urgent	hot	\
0	-0.067792	-0.002879	0.138664	-0.04772	-0.002571	-0.044136	
1	-0.067792	-0.002820	-0.011578	-0.04772	-0.002571	-0.044136	
2	-0.067792	-0.002824	0.014179	-0.04772	-0.002571	-0.044136	
3	-0.067792	-0.002840	0.014179	-0.04772	-0.002571	-0.044136	
4	-0.067792	-0.002842	0.035214	-0.04772	-0.002571	-0.044136	
5	-0.067792	-0.002842	0.035214	-0.04772	-0.002571	-0.044136	
6	-0.067792	-0.002847	0.032429	-0.04772	-0.002571	-0.044136	
7	-0.067792	-0.002901	0.097411	-0.04772	-0.002571	-0.044136	
8	-0.067792	-0.002849	-0.021717	-0.04772	-0.002571	-0.044136	
9	-0.067792	-0.002847	-0.002498	-0.04772	-0.002571	1.234470	

	num_failed_logins	num_compromised	num_root	num_file_creations	\
0	-0.009782	-0.005679	-0.00564	-0.011232	
1	-0.009782	-0.005679	-0.00564	-0.011232	
2	-0.009782	-0.005679	-0.00564	-0.011232	
3	-0.009782	-0.005679	-0.00564	-0.011232	
4	-0.009782	-0.005679	-0.00564	-0.011232	

5	-0.009782	-0.005679	-0.00564	-0.011232
6	-0.009782	-0.005679	-0.00564	-0.011232
7	-0.009782	-0.005679	-0.00564	-0.011232
8	-0.009782	-0.005679	-0.00564	-0.011232
9	-0.009782	-0.005679	-0.00564	-0.011232

	...	dst_host_serror_rate	dst_host_srv_serror_rate	\
0	...	-0.464418	-0.463202	
1	...	-0.464418	-0.463202	
2	...	-0.464418	-0.463202	
3	...	-0.464418	-0.463202	
4	...	-0.464418	-0.463202	
5	...	-0.464418	-0.463202	
6	...	-0.464418	-0.463202	
7	...	-0.464418	-0.463202	
8	...	-0.464418	-0.463202	
9	...	-0.464418	-0.463202	

	dst_host_rerror_rate	dst_host_srv_rerror_rate	land	logged_in	\
0	-0.25204	-0.249464	0	1	
1	-0.25204	-0.249464	0	1	
2	-0.25204	-0.249464	0	1	
3	-0.25204	-0.249464	0	1	
4	-0.25204	-0.249464	0	1	
5	-0.25204	-0.249464	0	1	
6	-0.25204	-0.249464	0	1	
7	-0.25204	-0.249464	0	1	
8	-0.25204	-0.249464	0	1	
9	-0.25204	-0.249464	0	1	

	root_shell	su_attempted	is_host_login	is_guest_login
0	0	0	0	0
1	0	0	0	0
2	0	0	0	0
3	0	0	0	0
4	0	0	0	0
5	0	0	0	0
6	0	0	0	0
7	0	0	0	0
8	0	0	0	0
9	0	0	0	0

[10 rows x 38 columns]

In [10]: *#Applying the same transformations on the test dataset*

```
kdd_test_10per_var = kdd_test_10per_mod.iloc[:,0:38]
kdd_test_10per_var.shape
```

```

kdd_test_10per_var.describe()
kdd_test_10per_var.head(10)
type(kdd_test_10per_var)

#Creating a list of the discrete valued columns (mentioned in the documentation for the
cont_labels = ["duration", "src_bytes", "dst_bytes", "wrong_fragment", "urgent", "hot", "num
               "num_compromised", "num_root", "num_file_creations", "num_shells", "num_access_files", "
               "count", "srv_count", "serror_rate", "srv_serror_rate", "rerror_rate", "srv_rerror_rate",
               "diff_srv_rate", "srv_diff_host_rate", "dst_host_count", "dst_host_srv_count", "dst_hos
               "dst_host_srv_diff_host_rate", "dst_host_serror_rate", "dst_host_srv_serror_rate", "ds
other_labels = ["land", "logged_in", "root_shell", "su_attempted", "is_host_login", "is_gues
kdd_test_10per_forscale = kdd_test_10per_mod[cont_labels]
kdd_test_10per_forscale.shape
kdd_test_10per_forscale.describe()

kdd_test_10per_others = kdd_test_10per_mod[other_labels]
kdd_test_10per_others.shape
kdd_test_10per_others.describe()

#scaling the test dataset
kdd_test_10per_scaled1 = pd.DataFrame(data = train_scaler.transform(kdd_test_10per_fors
kdd_test_10per_scaled1.shape
kdd_test_10per_scaled1.describe()
kdd_test_10per_scaled1.head(10)

#adding the binary valued variable columns
kdd_test_10per_scaled = pd.concat([kdd_test_10per_scaled1, kdd_test_10per_others], axis=
kdd_test_10per_scaled.shape
kdd_test_10per_scaled.describe()
kdd_test_10per_scaled.head(10)

```

Out[10]: (311029, 38)

```

Out[10]:
      duration  src_bytes  dst_bytes  land  \
count  311029.000000  3.110290e+05  3.110290e+05  311029.000000
mean      17.902736  1.731702e+03  7.479937e+02    0.000029
std      407.644400  1.276567e+05  1.612018e+04    0.005379
min         0.000000  0.000000e+00  0.000000e+00    0.000000
25%         0.000000  1.050000e+02  0.000000e+00    0.000000
50%         0.000000  5.200000e+02  0.000000e+00    0.000000
75%         0.000000  1.032000e+03  0.000000e+00    0.000000
max      57715.000000  6.282565e+07  5.203179e+06    1.000000

      wrong_fragment  urgent  hot  num_failed_logins  \
count  311029.000000  311029.000000  311029.000000  311029.000000
mean      0.000762    0.000051    0.014677    0.002363
std      0.040367    0.009821    0.312068    0.049990
min         0.000000    0.000000    0.000000    0.000000

```

25%	0.000000	0.000000	0.000000	0.000000
50%	0.000000	0.000000	0.000000	0.000000
75%	0.000000	0.000000	0.000000	0.000000
max	3.000000	3.000000	101.000000	4.000000

	logged_in	num_compromised	...	\
count	311029.000000	311029.000000	...	
mean	0.172476	0.011243	...	
std	0.377794	1.958325	...	
min	0.000000	0.000000	...	
25%	0.000000	0.000000	...	
50%	0.000000	0.000000	...	
75%	0.000000	0.000000	...	
max	1.000000	796.000000	...	

	dst_host_count	dst_host_srv_count	dst_host_same_srv_rate	\
count	311029.000000	311029.000000	311029.000000	
mean	235.282681	199.193914	0.793494	
std	60.913298	100.306470	0.387090	
min	0.000000	0.000000	0.000000	
25%	255.000000	244.000000	0.970000	
50%	255.000000	255.000000	1.000000	
75%	255.000000	255.000000	1.000000	
max	255.000000	255.000000	1.000000	

	dst_host_diff_srv_rate	dst_host_same_src_port_rate	\
count	311029.000000	311029.000000	
mean	0.024953	0.547919	
std	0.096003	0.491963	
min	0.000000	0.000000	
25%	0.000000	0.000000	
50%	0.000000	1.000000	
75%	0.010000	1.000000	
max	1.000000	1.000000	

	dst_host_srv_diff_host_rate	dst_host_serror_rate	\
count	311029.000000	311029.000000	
mean	0.004566	0.058764	
std	0.035773	0.231296	
min	0.000000	0.000000	
25%	0.000000	0.000000	
50%	0.000000	0.000000	
75%	0.000000	0.000000	
max	1.000000	1.000000	

	dst_host_srv_serror_rate	dst_host_rerror_rate	\
count	311029.000000	311029.000000	
mean	0.058791	0.142659	

std	0.232997	0.344380
min	0.000000	0.000000
25%	0.000000	0.000000
50%	0.000000	0.000000
75%	0.000000	0.000000
max	1.000000	1.000000

dst_host_srv_rerror_rate	
count	311029.000000
mean	0.141693
std	0.346573
min	0.000000
25%	0.000000
50%	0.000000
75%	0.000000
max	1.000000

[8 rows x 38 columns]

```
Out[10]:
```

	duration	src_bytes	dst_bytes	land	wrong_fragment	urgent	hot	\
0	0	105	146	0	0	0	0	
1	0	105	146	0	0	0	0	
2	0	105	146	0	0	0	0	
3	0	105	146	0	0	0	0	
4	0	105	146	0	0	0	0	
5	0	105	146	0	0	0	0	
6	0	29	0	0	0	0	0	
7	0	105	146	0	0	0	0	
8	0	105	146	0	0	0	0	
9	0	223	185	0	0	0	0	

	num_failed_logins	logged_in	num_compromised	...	\
0	0	0	0	...	
1	0	0	0	...	
2	0	0	0	...	
3	0	0	0	...	
4	0	0	0	...	
5	0	0	0	...	
6	0	0	0	...	
7	0	0	0	...	
8	0	0	0	...	
9	0	1	0	...	

	dst_host_count	dst_host_srv_count	dst_host_same_srv_rate	\
0	255	254	1.00	
1	255	254	1.00	
2	255	254	1.00	
3	255	254	1.00	

4	255	254	1.00
5	255	255	1.00
6	10	3	0.30
7	255	253	0.99
8	255	254	1.00
9	71	255	1.00

	dst_host_diff_srv_rate	dst_host_same_src_port_rate \
0	0.01	0.00
1	0.01	0.00
2	0.01	0.00
3	0.01	0.00
4	0.01	0.01
5	0.00	0.01
6	0.30	0.30
7	0.01	0.00
8	0.01	0.00
9	0.00	0.01

	dst_host_srv_diff_host_rate	dst_host_serror_rate \
0	0.00	0.0
1	0.00	0.0
2	0.00	0.0
3	0.00	0.0
4	0.00	0.0
5	0.00	0.0
6	0.00	0.0
7	0.00	0.0
8	0.00	0.0
9	0.01	0.0

	dst_host_srv_serror_rate	dst_host_rerror_rate	dst_host_srv_rerror_rate
0	0.0	0.0	0.0
1	0.0	0.0	0.0
2	0.0	0.0	0.0
3	0.0	0.0	0.0
4	0.0	0.0	0.0
5	0.0	0.0	0.0
6	0.0	0.0	0.0
7	0.0	0.0	0.0
8	0.0	0.0	0.0
9	0.0	0.0	0.0

[10 rows x 38 columns]

Out[10]: pandas.core.frame.DataFrame

Out[10]: (311029, 32)

```

Out[10]:
      duration      src_bytes      dst_bytes  wrong_fragment  \
count  311029.000000  3.110290e+05  3.110290e+05  311029.000000
mean    17.902736  1.731702e+03  7.479937e+02    0.000762
std    407.644400  1.276567e+05  1.612018e+04    0.040367
min      0.000000  0.000000e+00  0.000000e+00    0.000000
25%      0.000000  1.050000e+02  0.000000e+00    0.000000
50%      0.000000  5.200000e+02  0.000000e+00    0.000000
75%      0.000000  1.032000e+03  0.000000e+00    0.000000
max    57715.000000  6.282565e+07  5.203179e+06    3.000000

      urgent      hot  num_failed_logins  num_compromised  \
count  311029.000000  311029.000000  311029.000000  311029.000000
mean    0.000051    0.014677    0.002363    0.011243
std    0.009821    0.312068    0.049990    1.958325
min      0.000000    0.000000    0.000000    0.000000
25%      0.000000    0.000000    0.000000    0.000000
50%      0.000000    0.000000    0.000000    0.000000
75%      0.000000    0.000000    0.000000    0.000000
max      3.000000   101.000000    4.000000   796.000000

      num_root  num_file_creations  ...  \
count  311029.000000  311029.000000  ...
mean    0.008359    0.000958  ...
std     2.165196    0.193119  ...
min      0.000000    0.000000  ...
25%      0.000000    0.000000  ...
50%      0.000000    0.000000  ...
75%      0.000000    0.000000  ...
max      878.000000   100.000000  ...

      dst_host_count  dst_host_srv_count  dst_host_same_srv_rate  \
count  311029.000000  311029.000000  311029.000000
mean    235.282681    199.193914    0.793494
std     60.913298    100.306470    0.387090
min      0.000000    0.000000    0.000000
25%     255.000000    244.000000    0.970000
50%     255.000000    255.000000    1.000000
75%     255.000000    255.000000    1.000000
max     255.000000    255.000000    1.000000

      dst_host_diff_srv_rate  dst_host_same_src_port_rate  \
count  311029.000000  311029.000000
mean    0.024953    0.547919
std     0.096003    0.491963
min      0.000000    0.000000
25%      0.000000    0.000000
50%      0.000000    1.000000
75%      0.010000    1.000000

```

max	1.000000	1.000000
-----	----------	----------

	dst_host_srv_diff_host_rate	dst_host_serror_rate \
count	311029.000000	311029.000000
mean	0.004566	0.058764
std	0.035773	0.231296
min	0.000000	0.000000
25%	0.000000	0.000000
50%	0.000000	0.000000
75%	0.000000	0.000000
max	1.000000	1.000000

	dst_host_srv_serror_rate	dst_host_rerror_rate \
count	311029.000000	311029.000000
mean	0.058791	0.142659
std	0.232997	0.344380
min	0.000000	0.000000
25%	0.000000	0.000000
50%	0.000000	0.000000
75%	0.000000	0.000000
max	1.000000	1.000000

	dst_host_srv_rerror_rate
count	311029.000000
mean	0.141693
std	0.346573
min	0.000000
25%	0.000000
50%	0.000000
75%	0.000000
max	1.000000

[8 rows x 32 columns]

Out[10]: (311029, 6)

Out[10]:

	land	logged_in	root_shell	su_attempted \
count	311029.000000	311029.000000	311029.000000	311029.000000
mean	0.000029	0.172476	0.000199	0.000023
std	0.005379	0.377794	0.014117	0.005947
min	0.000000	0.000000	0.000000	0.000000
25%	0.000000	0.000000	0.000000	0.000000
50%	0.000000	0.000000	0.000000	0.000000
75%	0.000000	0.000000	0.000000	0.000000
max	1.000000	1.000000	1.000000	2.000000

	is_host_login	is_guest_login
count	311029.000000	311029.000000

mean	0.000039	0.002424
std	0.006211	0.049177
min	0.000000	0.000000
25%	0.000000	0.000000
50%	0.000000	0.000000
75%	0.000000	0.000000
max	1.000000	1.000000

Out[10]: (311029, 32)

Out[10]:

	duration	src_bytes	dst_bytes	wrong_fragment	\
count	311029.000000	311029.000000	311029.000000	311029.000000	
mean	-0.042496	-0.001309	-0.003648	-0.042068	
std	0.575976	0.129179	0.487899	0.299447	
min	-0.067792	-0.003062	-0.026287	-0.047720	
25%	-0.067792	-0.002955	-0.026287	-0.047720	
50%	-0.067792	-0.002535	-0.026287	-0.047720	
75%	-0.067792	-0.002017	-0.026287	-0.047720	
max	81.479854	63.571681	157.455071	22.206629	

	urgent	hot	num_failed_logins	num_compromised	\
count	311029.000000	311029.000000	311029.000000	311029.000000	
mean	0.006764	-0.025370	0.142485	0.000573	
std	1.782311	0.399013	3.221101	1.088972	
min	-0.002571	-0.044136	-0.009782	-0.005679	
25%	-0.002571	-0.044136	-0.009782	-0.005679	
50%	-0.002571	-0.044136	-0.009782	-0.005679	
75%	-0.002571	-0.044136	-0.009782	-0.005679	
max	544.437104	129.095065	257.729132	442.628578	

	num_root	num_file_creations	...	\
count	311029.000000	311029.000000	...	
mean	-0.001487	-0.001295	...	
std	1.075758	2.002976	...	
min	-0.005640	-0.011232	...	
25%	-0.005640	-0.011232	...	
50%	-0.005640	-0.011232	...	
75%	-0.005640	-0.011232	...	
max	436.220770	1037.163376	...	

	dst_host_count	dst_host_srv_count	dst_host_same_srv_rate	\
count	311029.000000	311029.000000	311029.000000	
mean	0.043430	0.099285	0.096681	
std	0.940814	0.945928	0.942328	
min	-3.590542	-1.779188	-1.834994	
25%	0.347967	0.521823	0.526364	
50%	0.347967	0.625558	0.599396	
75%	0.347967	0.625558	0.599396	

max	0.347967	0.625558	0.599396
-----	----------	----------	----------

	dst_host_diff_srv_rate	dst_host_same_src_port_rate	\
count	311029.000000	311029.000000	
mean	-0.054486	-0.112228	
std	0.878676	1.022135	
min	-0.282867	-1.250621	
25%	-0.282867	-1.250621	
50%	-0.282867	0.827048	
75%	-0.191341	0.827048	
max	8.869697	0.827048	

	dst_host_srv_diff_host_rate	dst_host_serror_rate	\
count	311029.000000	311029.000000	
mean	-0.050252	-0.310016	
std	0.849060	0.607725	
min	-0.158629	-0.464418	
25%	-0.158629	-0.464418	
50%	-0.158629	-0.464418	
75%	-0.158629	-0.464418	
max	23.575830	2.163063	

	dst_host_srv_serror_rate	dst_host_rerror_rate	\
count	311029.000000	311029.000000	
mean	-0.308862	0.366631	
std	0.611671	1.493477	
min	-0.463202	-0.252040	
25%	-0.463202	-0.252040	
50%	-0.463202	-0.252040	
75%	-0.463202	-0.252040	
max	2.162027	4.084676	

	dst_host_srv_rerror_rate
count	311029.000000
mean	0.366219
std	1.505923
min	-0.249464
25%	-0.249464
50%	-0.249464
75%	-0.249464
max	4.095715

[8 rows x 32 columns]

```
Out[10]:
```

	duration	src_bytes	dst_bytes	wrong_fragment	urgent	hot	\
0	-0.067792	-0.002955	-0.021868	-0.04772	-0.002571	-0.044136	
1	-0.067792	-0.002955	-0.021868	-0.04772	-0.002571	-0.044136	
2	-0.067792	-0.002955	-0.021868	-0.04772	-0.002571	-0.044136	

3	-0.067792	-0.002955	-0.021868	-0.04772	-0.002571	-0.044136
4	-0.067792	-0.002955	-0.021868	-0.04772	-0.002571	-0.044136
5	-0.067792	-0.002955	-0.021868	-0.04772	-0.002571	-0.044136
6	-0.067792	-0.003032	-0.026287	-0.04772	-0.002571	-0.044136
7	-0.067792	-0.002955	-0.021868	-0.04772	-0.002571	-0.044136
8	-0.067792	-0.002955	-0.021868	-0.04772	-0.002571	-0.044136
9	-0.067792	-0.002836	-0.020688	-0.04772	-0.002571	-0.044136

	num_failed_logins	num_compromised	num_root	num_file_creations	\
0	-0.009782	-0.005679	-0.00564	-0.011232	
1	-0.009782	-0.005679	-0.00564	-0.011232	
2	-0.009782	-0.005679	-0.00564	-0.011232	
3	-0.009782	-0.005679	-0.00564	-0.011232	
4	-0.009782	-0.005679	-0.00564	-0.011232	
5	-0.009782	-0.005679	-0.00564	-0.011232	
6	-0.009782	-0.005679	-0.00564	-0.011232	
7	-0.009782	-0.005679	-0.00564	-0.011232	
8	-0.009782	-0.005679	-0.00564	-0.011232	
9	-0.009782	-0.005679	-0.00564	-0.011232	

	...	dst_host_count	dst_host_srv_count	\
0	...	0.347967	0.616127	
1	...	0.347967	0.616127	
2	...	0.347967	0.616127	
3	...	0.347967	0.616127	
4	...	0.347967	0.616127	
5	...	0.347967	0.625558	
6	...	-3.436091	-1.750897	
7	...	0.347967	0.606697	
8	...	0.347967	0.616127	
9	...	-2.493938	0.625558	

	dst_host_same_srv_rate	dst_host_diff_srv_rate	\
0	0.599396	-0.191341	
1	0.599396	-0.191341	
2	0.599396	-0.191341	
3	0.599396	-0.191341	
4	0.599396	-0.191341	
5	0.599396	-0.282867	
6	-1.104677	2.462903	
7	0.575052	-0.191341	
8	0.599396	-0.191341	
9	0.599396	-0.282867	

	dst_host_same_src_port_rate	dst_host_srv_diff_host_rate	\
0	-1.250621	-0.158629	
1	-1.250621	-0.158629	
2	-1.250621	-0.158629	

3	-1.250621	-0.158629
4	-1.229844	-0.158629
5	-1.229844	-0.158629
6	-0.627320	-0.158629
7	-1.250621	-0.158629
8	-1.250621	-0.158629
9	-1.229844	0.078715

	dst_host_serror_rate	dst_host_srv_serror_rate	dst_host_rerror_rate	\
0	-0.464418	-0.463202	-0.25204	
1	-0.464418	-0.463202	-0.25204	
2	-0.464418	-0.463202	-0.25204	
3	-0.464418	-0.463202	-0.25204	
4	-0.464418	-0.463202	-0.25204	
5	-0.464418	-0.463202	-0.25204	
6	-0.464418	-0.463202	-0.25204	
7	-0.464418	-0.463202	-0.25204	
8	-0.464418	-0.463202	-0.25204	
9	-0.464418	-0.463202	-0.25204	

	dst_host_srv_rerror_rate
0	-0.249464
1	-0.249464
2	-0.249464
3	-0.249464
4	-0.249464
5	-0.249464
6	-0.249464
7	-0.249464
8	-0.249464
9	-0.249464

[10 rows x 32 columns]

Out[10]: (311029, 38)

Out[10]:

	duration	src_bytes	dst_bytes	wrong_fragment	\
count	311029.000000	311029.000000	311029.000000	311029.000000	
mean	-0.042496	-0.001309	-0.003648	-0.042068	
std	0.575976	0.129179	0.487899	0.299447	
min	-0.067792	-0.003062	-0.026287	-0.047720	
25%	-0.067792	-0.002955	-0.026287	-0.047720	
50%	-0.067792	-0.002535	-0.026287	-0.047720	
75%	-0.067792	-0.002017	-0.026287	-0.047720	
max	81.479854	63.571681	157.455071	22.206629	

	urgent	hot	num_failed_logins	num_compromised	\
count	311029.000000	311029.000000	311029.000000	311029.000000	

mean	0.006764	-0.025370	0.142485	0.000573
std	1.782311	0.399013	3.221101	1.088972
min	-0.002571	-0.044136	-0.009782	-0.005679
25%	-0.002571	-0.044136	-0.009782	-0.005679
50%	-0.002571	-0.044136	-0.009782	-0.005679
75%	-0.002571	-0.044136	-0.009782	-0.005679
max	544.437104	129.095065	257.729132	442.628578

	num_root	num_file_creations	...	\
count	311029.000000	311029.000000	...	
mean	-0.001487	-0.001295	...	
std	1.075758	2.002976	...	
min	-0.005640	-0.011232	...	
25%	-0.005640	-0.011232	...	
50%	-0.005640	-0.011232	...	
75%	-0.005640	-0.011232	...	
max	436.220770	1037.163376	...	

	dst_host_serror_rate	dst_host_srv_serror_rate	dst_host_rerror_rate	\
count	311029.000000	311029.000000	311029.000000	
mean	-0.310016	-0.308862	0.366631	
std	0.607725	0.611671	1.493477	
min	-0.464418	-0.463202	-0.252040	
25%	-0.464418	-0.463202	-0.252040	
50%	-0.464418	-0.463202	-0.252040	
75%	-0.464418	-0.463202	-0.252040	
max	2.163063	2.162027	4.084676	

	dst_host_srv_rerror_rate	land	logged_in	root_shell	\
count	311029.000000	311029.000000	311029.000000	311029.000000	
mean	0.366219	0.000029	0.172476	0.000199	
std	1.505923	0.005379	0.377794	0.014117	
min	-0.249464	0.000000	0.000000	0.000000	
25%	-0.249464	0.000000	0.000000	0.000000	
50%	-0.249464	0.000000	0.000000	0.000000	
75%	-0.249464	0.000000	0.000000	0.000000	
max	4.095715	1.000000	1.000000	1.000000	

	su_attempted	is_host_login	is_guest_login
count	311029.000000	311029.000000	311029.000000
mean	0.000023	0.000039	0.002424
std	0.005947	0.006211	0.049177
min	0.000000	0.000000	0.000000
25%	0.000000	0.000000	0.000000
50%	0.000000	0.000000	0.000000
75%	0.000000	0.000000	0.000000
max	2.000000	1.000000	1.000000

[8 rows x 38 columns]

```
Out[10]:  duration  src_bytes  dst_bytes  wrong_fragment  urgent  hot  \
0 -0.067792 -0.002955 -0.021868 -0.04772 -0.002571 -0.044136
1 -0.067792 -0.002955 -0.021868 -0.04772 -0.002571 -0.044136
2 -0.067792 -0.002955 -0.021868 -0.04772 -0.002571 -0.044136
3 -0.067792 -0.002955 -0.021868 -0.04772 -0.002571 -0.044136
4 -0.067792 -0.002955 -0.021868 -0.04772 -0.002571 -0.044136
5 -0.067792 -0.002955 -0.021868 -0.04772 -0.002571 -0.044136
6 -0.067792 -0.003032 -0.026287 -0.04772 -0.002571 -0.044136
7 -0.067792 -0.002955 -0.021868 -0.04772 -0.002571 -0.044136
8 -0.067792 -0.002955 -0.021868 -0.04772 -0.002571 -0.044136
9 -0.067792 -0.002836 -0.020688 -0.04772 -0.002571 -0.044136

    num_failed_logins  num_compromised  num_root  num_file_creations  \
0 -0.009782 -0.005679 -0.00564 -0.011232
1 -0.009782 -0.005679 -0.00564 -0.011232
2 -0.009782 -0.005679 -0.00564 -0.011232
3 -0.009782 -0.005679 -0.00564 -0.011232
4 -0.009782 -0.005679 -0.00564 -0.011232
5 -0.009782 -0.005679 -0.00564 -0.011232
6 -0.009782 -0.005679 -0.00564 -0.011232
7 -0.009782 -0.005679 -0.00564 -0.011232
8 -0.009782 -0.005679 -0.00564 -0.011232
9 -0.009782 -0.005679 -0.00564 -0.011232

    ...  dst_host_serror_rate  dst_host_srv_serror_rate  \
0 ... -0.464418 -0.463202
1 ... -0.464418 -0.463202
2 ... -0.464418 -0.463202
3 ... -0.464418 -0.463202
4 ... -0.464418 -0.463202
5 ... -0.464418 -0.463202
6 ... -0.464418 -0.463202
7 ... -0.464418 -0.463202
8 ... -0.464418 -0.463202
9 ... -0.464418 -0.463202

    dst_host_rerror_rate  dst_host_srv_rerror_rate  land  logged_in  \
0 -0.25204 -0.249464 0 0
1 -0.25204 -0.249464 0 0
2 -0.25204 -0.249464 0 0
3 -0.25204 -0.249464 0 0
4 -0.25204 -0.249464 0 0
5 -0.25204 -0.249464 0 0
6 -0.25204 -0.249464 0 0
7 -0.25204 -0.249464 0 0
8 -0.25204 -0.249464 0 0
```

9	-0.25204	-0.249464	0	1
	root_shell	su_attempted	is_host_login	is_guest_login
0	0	0	0	0
1	0	0	0	0
2	0	0	0	0
3	0	0	0	0
4	0	0	0	0
5	0	0	0	0
6	0	0	0	0
7	0	0	0	0
8	0	0	0	0
9	0	0	0	0

[10 rows x 38 columns]

1.1.4 Data Analysis

Now we use the final datasets to train and run models based on the following algorithms and evaluate the best one for detection and prediction purposes: - Logistic Regression - K-Nearest Neighbors - Support Vector Machine - Random Forest - Neural Networks - Ensemble Modeling - Unsupervised Learning - Clustering

Since there are several different attack types mentioned in the "attack_type" label, we first classify them into "normal" and "bad" (as done in the Data Exploration section) and try to develop a model that detects "bad" logs (majority of them being DoS attacks) in the test dataset. We also implement multi-class detection for these models as well.

```
In [11]: #Importing the required packages

from sklearn import metrics
from sklearn.linear_model import LogisticRegression
from sklearn.neighbors import KNeighborsClassifier
from scipy import stats
import pylab as pl
from sklearn.svm import SVC
from sklearn.ensemble import RandomForestClassifier
from sklearn.cluster import KMeans
from sklearn.ensemble import BaggingClassifier
from sklearn.neural_network import MLPClassifier
from sklearn.ensemble import VotingClassifier
from sklearn import tree
```

1.1.5 Logistic Regression

We use a logistic regression model trained on the training dataset to predict "normal" and "bad" as well as multi-class labels for the test dataset.

```
In [14]: #Creating the "normal" and "bad" labels for test dataset as well
attack_labs_test = kdd_test_10per_mod['attack_type'].copy()
```

```

attack_labs_test[attack_labs_test!='normal.'] = 'bad.'
attack_labs_test.value_counts()

#Logistic Regression Model -- 2 labels
logrmodel = LogisticRegression(C=0.01)
logrmodel.fit(kdd_train_10per_scaled, attack_labs)
log_preds = logrmodel.predict(kdd_test_10per_scaled)
print(log_preds)

#Logistic Regression Model -- multi-labels
logrmodel1 = LogisticRegression(C=0.01,solver='sag',multi_class='ovr')
logrmodel1.fit(kdd_train_10per_scaled, attack_labs_all)
log_preds1 = logrmodel1.predict(kdd_test_10per_scaled)
print(log_preds1)

#Calculating test metrics to evaluate the model -- 2 labels
#Confusion Matrix, accuracy, recall, precision and f1-score
print(metrics.confusion_matrix(attack_labs_test, log_preds))
print(metrics.accuracy_score(attack_labs_test, log_preds))
print(metrics.recall_score(attack_labs_test, log_preds,average='macro'))
print(metrics.precision_score(attack_labs_test, log_preds,average='macro'))
print(metrics.f1_score(attack_labs_test, log_preds,average='macro'))

ht = sns.heatmap(pd.DataFrame(metrics.confusion_matrix(attack_labs_test, log_preds)), a
ht.set(xlabel='Predicted Values', ylabel='True Values')
plt.show()

#Another way to view these metrics for every class label
precision, recall, fscore, support = metrics.precision_recall_fscore_support(attack_lab

print('precision: {}'.format(precision))
print('recall: {}'.format(recall))
print('fscore: {}'.format(fscore))
print('support: {}'.format(support))

#Calculating test metrics to evaluate the model -- multi-labels
#Confusion Matrix, accuracy, recall, precision and f1-score
print(metrics.confusion_matrix(attack_labs_tall, log_preds1))
print(metrics.accuracy_score(attack_labs_tall, log_preds1))
print(metrics.recall_score(attack_labs_tall, log_preds1,average='macro'))
print(metrics.precision_score(attack_labs_tall, log_preds1,average='macro'))
print(metrics.f1_score(attack_labs_tall, log_preds1,average='macro'))

ht = sns.heatmap(pd.DataFrame(metrics.confusion_matrix(attack_labs_tall, log_preds1)),
ht.set(xlabel='Predicted Values', ylabel='True Values')
plt.show()

```



```

#Another way to view these metrics for every class label
precision, recall, fscore, support = metrics.precision_recall_fscore_support(attack_label,
                                     predicted_label)

print('precision: {}'.format(precision))
print('recall: {}'.format(recall))
print('fscore: {}'.format(fscore))
print('support: {}'.format(support))

```

```

Out[14]: bad.          250436
        normal.       60593
        Name: attack_type, dtype: int64

```

```

Out[14]: LogisticRegression(C=0.01, class_weight=None, dual=False, fit_intercept=True,
                             intercept_scaling=1, max_iter=100, multi_class='ovr', n_jobs=1,
                             penalty='l2', random_state=None, solver='liblinear', tol=0.0001,
                             verbose=0, warm_start=False)

```

```

['normal.' 'normal.' 'normal.' ..., 'normal.' 'normal.' 'normal.']

```

```

C:\Users\Neha Rawat\Anaconda3\lib\site-packages\sklearn\linear_model\sag.py:326: ConvergenceWarning:
  "the coef_ did not converge", ConvergenceWarning)

```

```

Out[14]: LogisticRegression(C=0.01, class_weight=None, dual=False, fit_intercept=True,
                             intercept_scaling=1, max_iter=100, multi_class='ovr', n_jobs=1,
                             penalty='l2', random_state=None, solver='sag', tol=0.0001,
                             verbose=0, warm_start=False)

```

```

['normal.' 'normal.' 'normal.' ..., 'normal.' 'normal.' 'normal.']

```

```


[[225627  24809]
 [   994  59599]]
0.917039890171
0.942266115551
0.850848092029
0.883980813408

```

```


Out[14]: [Text(33,0.5,'True Values'), Text(0.5,15,'Predicted Values')]

```



```
precision: [ 0.99561382  0.70608236]
recall:   [ 0.90093677  0.98359546]
fscore:   [ 0.94591212  0.8220495 ]
support:  [250436  60593]
[[222373      0      0  7446    34]
 [      6    23      0 16148    12]
 [     49     1      6   169     3]
 [    790    26      1 59699    77]
 [    200     0      0  1341   2625]]
0.915432323031
0.522108121159
0.794127347192
0.523056025877
```

```
Out[14]: [Text(33,0.5,'True Values'), Text(0.5,15,'Predicted Values')]
```



```
precision: [ 0.99532267  0.46          0.85714286  0.70397274  0.95419847]
recall:    [ 0.96745746  0.00142072  0.02631579  0.98524582  0.63010082]
fscore:    [ 0.98119227  0.00283269  0.05106383  0.82119178  0.75899957]
support:   [229853  16189    228  60593   4166]
```

We see that the accuracy (91.70 percent) and F1 scores (88.39 percent) for 2-label predictions are good enough. Through the precision_recall_fscore_support matrix we can also see that correct prediction of "bad" labels is higher, due to higher proportion of these labels in the test as well as training data. When we try to fit a multi-class problem using logistic regression, we observe that the recall for "R2L" and "U2R" labels is very less (due to small proportion of these labels in the training dataset), whereas predictions for DOS and "normal" labels is quite good. Therefore, we will proceed to try other algorithms to try and achieve better and more balanced prediction scores.

1.1.6 K-Nearest Neighbors

We use a K-Nearest Neighbors model (with k=5 and Manhattan Distance Metric) trained on the training dataset to predict "normal" and "bad" as well as multi-class labels for the test dataset.

```
In [15]: #K-Nearest Neighbors
         #We run the K-nearest neighbors model for k = 5
         #Takes a lot of time -- around two hours

         #K=5 -- 2 labels
         knn_mod1 = KNeighborsClassifier(n_neighbors = 5, algorithm = 'ball_tree', leaf_size=600)
```

```

knn_mod1.fit(kdd_train_10per_scaled, attack_labs)
knn_preds1 = knn_mod1.predict(kdd_test_10per_scaled)
print(knn_preds1)

#Multi-labels
knn_mod2 = KNeighborsClassifier(n_neighbors = 5, algorithm = 'ball_tree', leaf_size=600)
knn_mod2.fit(kdd_train_10per_scaled, attack_labs_all)
knn_preds2 = knn_mod2.predict(kdd_test_10per_scaled)
print(knn_preds2)

Out[15]: KNeighborsClassifier(algorithm='ball_tree', leaf_size=600, metric='manhattan',
                             metric_params=None, n_jobs=-1, n_neighbors=5, p=2,
                             weights='uniform')

['normal.' 'normal.' 'normal.' ..., 'normal.' 'normal.' 'normal.']

Out[15]: KNeighborsClassifier(algorithm='ball_tree', leaf_size=600, metric='manhattan',
                             metric_params=None, n_jobs=-1, n_neighbors=5, p=2,
                             weights='uniform')

['normal.' 'normal.' 'normal.' ..., 'normal.' 'normal.' 'normal.']

In [16]: #Evaluation Metrics for KNN k=5 -- 2 labels
         #Confusion Matrix, accuracy, recall, precision and f1-score
print(metrics.confusion_matrix(attack_labs_test, knn_preds1))
print(metrics.accuracy_score(attack_labs_test, knn_preds1))
print(metrics.recall_score(attack_labs_test, knn_preds1, average='macro'))
print(metrics.precision_score(attack_labs_test, knn_preds1, average='macro'))
print(metrics.f1_score(attack_labs_test, knn_preds1, average='macro'))

ht = sns.heatmap(pd.DataFrame(metrics.confusion_matrix(attack_labs_test, knn_preds1)),
ht.set(xlabel='Predicted Values', ylabel='True Values')
plt.show()

#Another way to view these metrics for every class label
precision, recall, fscore, support = metrics.precision_recall_fscore_support(attack_labs_test, knn_preds1)

print('precision: {}'.format(precision))
print('recall: {}'.format(recall))
print('fscore: {}'.format(fscore))
print('support: {}'.format(support))

#Evaluation Metrics for KNN k=5 -- multi-labels
#Confusion Matrix, accuracy, recall, precision and f1-score
print(metrics.confusion_matrix(attack_labs_tall, knn_preds2))
print(metrics.accuracy_score(attack_labs_tall, knn_preds2))
print(metrics.recall_score(attack_labs_tall, knn_preds2, average='macro'))

```

```

print(metrics.precision_score(attack_labs_tall, knn_preds2, average='macro'))
print(metrics.f1_score(attack_labs_tall, knn_preds2, average='macro'))

ht = sns.heatmap(pd.DataFrame(metrics.confusion_matrix(attack_labs_tall, knn_preds2)),
ht.set(xlabel='Predicted Values', ylabel='True Values')
plt.show()

#Another way to view these metrics for every class label
precision, recall, fscore, support = metrics.precision_recall_fscore_support(attack_lab

print('precision: {}'.format(precision))
print('recall: {}'.format(recall))
print('fscore: {}'.format(fscore))
print('support: {}'.format(support))


```

```

[[227642  22794]
 [   977 59616]]
0.923573043028
0.946429379698
0.859566933948
0.892076055409

```

Out[16]: [Text(33,0.5,'True Values'), Text(0.5,15,'Predicted Values')]

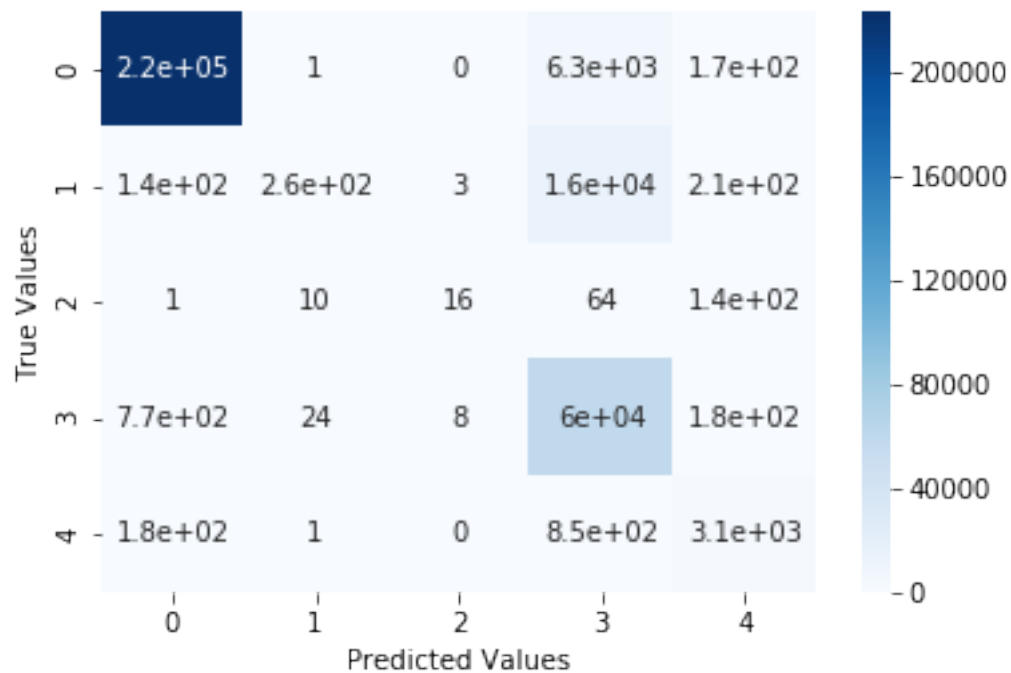


```

precision: [ 0.99572651  0.72340735]
recall: [ 0.90898273  0.98387603]
fscore: [ 0.95037939  0.83377272]
support: [250436  60593]
[[223384      1      0  6301    167]
 [   135    258      3 15582    211]
 [      1     10     16     64    137]
 [   768     24      8 59617    176]
 [   178      1      0   849   3138]]
0.920856254561
0.55902023168
0.801650312904
0.551786905652

```

Out[16]: [Text(33,0.5,'True Values'), Text(0.5,15,'Predicted Values')]



```

precision: [ 0.99517967  0.87755102  0.59259259  0.72339315  0.81953513]
recall: [ 0.97185593  0.01593675  0.07017544  0.98389253  0.75324052]
fscore: [ 0.98337952  0.03130498  0.1254902  0.83376921  0.78499062]
support: [229853  16189    228  60593   4166]

```

We see that with the use of KNN (k=5) and Manhattan Distance metric we are able to improve the accuracy (92.357 percent) and F1 score (89.20 percent) slightly more than in logistic regression

for a 2-label classification. For multi-label classification as well, we see an improvement in recall scores overall and especially for "R2L" and "U2R" categories (as compared to logistic regression). Also, we see that the F1 score for "normal" labels has increased too.

1.1.7 Support Vector Machine

We use a Support Vector Machine (using a linear and polynomial kernel) model trained on the training dataset to predict "normal" and "bad" as well as multi-class labels on the test dataset.

```
In [17]: #Training using an SVM model (kernel = linear) -- 2 labels
```

```
svm_mod1 = BaggingClassifier(SVC(kernel='linear'),max_samples = 0.1, n_jobs=-1)
svm_mod1.fit(kdd_train_10per_scaled, attack_labs)
svm_preds1 = svm_mod1.predict(kdd_test_10per_scaled)
print(svm_preds1)

#Training using an SVM model (kernel = linear) -- multi-labels
svm_mod2 = BaggingClassifier(SVC(kernel='linear'),max_samples = 0.1, n_jobs=-1)
svm_mod2.fit(kdd_train_10per_scaled, attack_labs_all)
svm_preds2 = svm_mod2.predict(kdd_test_10per_scaled)
print(svm_preds2)
```

```
Out[17]: BaggingClassifier(base_estimator=SVC(C=1.0, cache_size=200, class_weight=None, coef0=0.0,
      decision_function_shape='ovr', degree=3, gamma='auto', kernel='linear',
      max_iter=-1, probability=False, random_state=None, shrinking=True,
      tol=0.001, verbose=False),
      bootstrap=True, bootstrap_features=False, max_features=1.0,
      max_samples=0.1, n_estimators=10, n_jobs=-1, oob_score=False,
      random_state=None, verbose=0, warm_start=False)
```

```
['normal.' 'normal.' 'normal.' ..., 'normal.' 'normal.' 'normal.']
```

```
Out[17]: BaggingClassifier(base_estimator=SVC(C=1.0, cache_size=200, class_weight=None, coef0=0.0,
      decision_function_shape='ovr', degree=3, gamma='auto', kernel='linear',
      max_iter=-1, probability=False, random_state=None, shrinking=True,
      tol=0.001, verbose=False),
      bootstrap=True, bootstrap_features=False, max_features=1.0,
      max_samples=0.1, n_estimators=10, n_jobs=-1, oob_score=False,
      random_state=None, verbose=0, warm_start=False)
```

```
['normal.' 'normal.' 'normal.' ..., 'normal.' 'normal.' 'normal.']
```

```
In [18]: #Evaluation Metrics for SVM (kernel = linear) -- 2 labels
```

```
#Confusion Matrix, accuracy, recall, precision and f1-score
print(metrics.confusion_matrix(attack_labs_test, svm_preds1))
print(metrics.accuracy_score(attack_labs_test, svm_preds1))
print(metrics.recall_score(attack_labs_test, svm_preds1,average='macro'))
```

```

print(metrics.precision_score(attack_labs_test, svm_preds1, average='macro'))
print(metrics.f1_score(attack_labs_test, svm_preds1, average='macro'))

ht = sns.heatmap(pd.DataFrame(metrics.confusion_matrix(attack_labs_test, svm_preds1)),
ht.set(xlabel='Predicted Values', ylabel='True Values')
plt.show()

#Another way to view these metrics for every class label
precision, recall, fscore, support = metrics.precision_recall_fscore_support(attack_labs_test, svm_preds1)

print('precision: {}'.format(precision))
print('recall: {}'.format(recall))
print('fscore: {}'.format(fscore))
print('support: {}'.format(support))

#Evaluation Metrics for SVM (kernel = linear) -- multi-labels
#Confusion Matrix, accuracy, recall, precision and f1-score
print(metrics.confusion_matrix(attack_labs_tall, svm_preds2))
print(metrics.accuracy_score(attack_labs_tall, svm_preds2))
print(metrics.recall_score(attack_labs_tall, svm_preds2, average='macro'))
print(metrics.precision_score(attack_labs_tall, svm_preds2, average='macro'))
print(metrics.f1_score(attack_labs_tall, svm_preds2, average='macro'))

ht = sns.heatmap(pd.DataFrame(metrics.confusion_matrix(attack_labs_tall, svm_preds2)),
ht.set(xlabel='Predicted Values', ylabel='True Values')
plt.show()

#Another way to view these metrics for every class label
precision, recall, fscore, support = metrics.precision_recall_fscore_support(attack_labs_tall, svm_preds2)

print('precision: {}'.format(precision))
print('recall: {}'.format(recall))
print('fscore: {}'.format(fscore))
print('support: {}'.format(support))

```

```


[[227261  23175]
 [  1003 59590]]
0.922264483376
0.945454160077
0.857798149551
0.890418839275

```

```


Out[18]: [Text(33,0.5,'True Values'), Text(0.5,15,'Predicted Values')]

```

```
precision: [ 0.99560597  0.71999033]
recall: [ 0.90746139  0.98344693]
fscore: [ 0.94949238  0.8313453 ]
support: [250436  60593]
[[213998      0      0 15852      3]
 [      9    718      3 15449    10]
 [     35      2     21    122    48]
 [    794     56      8 59638    97]
 [     229      0      0   1046 2891]]
0.891447421302
0.549133526469
0.834442449644
0.558129727845
```

```
Out[18]: [Text(33,0.5,'True Values'), Text(0.5,15,'Predicted Values')]
```



```
precision: [ 0.99503871  0.92525773  0.65625      0.64748608  0.94817973]
recall:    [ 0.93102113  0.0443511   0.09210526  0.9842391   0.69395103]
fscore:    [ 0.96196603  0.08464486  0.16153846  0.78111329  0.801386  ]
support:   [229853  16189    228  60593   4166]
```

```
In [19]: #SVM using a polynomial kernel -- 2 labels
svm_mod3 = BaggingClassifier(SVC(kernel='poly',cache_size=7000),max_samples = 0.1, n_jobs=10)
svm_mod3.fit(kdd_train_10per_scaled, attack_labs)
svm_preds3 = svm_mod3.predict(kdd_test_10per_scaled)
print(svm_preds3)

#SVM using a polynomial kernel -- multi-labels
svm_mod4 = BaggingClassifier(SVC(kernel='poly',cache_size=7000),max_samples = 0.1, n_jobs=10)
svm_mod4.fit(kdd_train_10per_scaled, attack_labs_all)
svm_preds4 = svm_mod4.predict(kdd_test_10per_scaled)
print(svm_preds4)
```

```
Out[19]: BaggingClassifier(base_estimator=SVC(C=1.0, cache_size=7000, class_weight=None, coef0=0.0,
decision_function_shape='ovr', degree=3, gamma='auto', kernel='poly',
max_iter=-1, probability=False, random_state=None, shrinking=True,
tol=0.001, verbose=False),
bootstrap=True, bootstrap_features=False, max_features=1.0,
max_samples=0.1, n_estimators=10, n_jobs=-1, oob_score=False,
random_state=None, verbose=0, warm_start=False)
```

```
['normal.' 'normal.' 'normal.' ..., 'normal.' 'normal.' 'normal.']
```

```
Out[19]: BaggingClassifier(base_estimator=SVC(C=1.0, cache_size=7000, class_weight=None, coef0=0.0,
      decision_function_shape='ovr', degree=3, gamma='auto', kernel='poly',
      max_iter=-1, probability=False, random_state=None, shrinking=True,
      tol=0.001, verbose=False),
      bootstrap=True, bootstrap_features=False, max_features=1.0,
      max_samples=0.1, n_estimators=10, n_jobs=-1, oob_score=False,
      random_state=None, verbose=0, warm_start=False)
```

```
['normal.' 'normal.' 'normal.' ..., 'normal.' 'normal.' 'normal.']
```

```
In [20]: #Evaluation Metrics for SVM (kernel = polynomial) -- 2 labels
#Confusion Matrix, accuracy, recall, precision and f1-score
print(metrics.confusion_matrix(attack_labs_test, svm_preds3))
print(metrics.accuracy_score(attack_labs_test, svm_preds3))
print(metrics.recall_score(attack_labs_test, svm_preds3, average='macro'))
print(metrics.precision_score(attack_labs_test, svm_preds3, average='macro'))
print(metrics.f1_score(attack_labs_test, svm_preds3, average='macro'))

ht = sns.heatmap(pd.DataFrame(metrics.confusion_matrix(attack_labs_test, svm_preds3)),
ht.set(xlabel='Predicted Values', ylabel='True Values')
plt.show()

#Another way to view these metrics for every class label
precision, recall, fscore, support = metrics.precision_recall_fscore_support(attack_labs_test, svm_preds3)

print('precision: {}'.format(precision))
print('recall: {}'.format(recall))
print('fscore: {}'.format(fscore))
print('support: {}'.format(support))

#Evaluation Metrics for SVM (kernel = polynomial) -- multi-labels
#Confusion Matrix, accuracy, recall, precision and f1-score
print(metrics.confusion_matrix(attack_labs_tall, svm_preds4))
print(metrics.accuracy_score(attack_labs_tall, svm_preds4))
print(metrics.recall_score(attack_labs_tall, svm_preds4, average='macro'))
print(metrics.precision_score(attack_labs_tall, svm_preds4, average='macro'))
print(metrics.f1_score(attack_labs_tall, svm_preds4, average='macro'))

ht = sns.heatmap(pd.DataFrame(metrics.confusion_matrix(attack_labs_tall, svm_preds4)),
ht.set(xlabel='Predicted Values', ylabel='True Values')
plt.show()

#Another way to view these metrics for every class label
precision, recall, fscore, support = metrics.precision_recall_fscore_support(attack_labs_tall, svm_preds4)
```

```

print('precision: {}'.format(precision))
print('recall: {}'.format(recall))
print('fscore: {}'.format(fscore))
print('support: {}'.format(support))


```

```

[[227707  22729]
 [  1274 59319]]
0.922827131875
0.94410837523
0.858707725634
0.890828669199

```

Out[20]: [Text(33,0.5,'True Values'), Text(0.5,15,'Predicted Values')]



```

precision: [ 0.99443622  0.72297923]
recall: [ 0.90924228  0.97897447]
fscore: [ 0.94993294  0.8317244 ]
support: [250436  60593]
[[222993    0    0   6723   137]
 [   12   354    1  15660   162]
 [    7    4   19    78   120]
 [  1049   33    5  59357   149]
 [    94    0    0    815  3257]]

```

```


0.919464101418
0.567352312799
0.846002193356
0.563851067765

```

```

Out[20]: [Text(33,0.5,'True Values'), Text(0.5,15,'Predicted Values')]

```



```

precision: [ 0.99481609  0.90537084  0.76          0.71832077  0.85150327]
recall:    [ 0.97015484  0.0218667   0.08333333  0.9796016   0.78180509]
fscore:    [ 0.98233071  0.04270205  0.15019763  0.82885789  0.81516706]
support:   [229853  16189    228  60593   4166]

```

We see that the polynomial SVM performs better than logistic regression but almost similar to KNN.

1.1.8 Random Forest

We now use a Random Forest model trained on the training dataset to predict "normal" and "bad" as well as multi-class labels on the test dataset.

```

In [21]: #Random Forest Model -- 2 labels
rdforestmod = RandomForestClassifier(n_estimators=10,n_jobs=-1, random_state=0)
rdforestmod.fit(kdd_train_10per_scaled, attack_labs)

```

```
rdpreds = rforestmod.predict(kdd_test_10per_scaled)
print(rdpreds)

#Random Forest Model -- multi-labels
rforestmod1 = RandomForestClassifier(n_estimators=10,n_jobs=-1, random_state=0)
rforestmod1.fit(kdd_train_10per_scaled, attack_labs_all)
rdpreds1 = rforestmod1.predict(kdd_test_10per_scaled)
print(rdpreds1)
```

```
Out[21]: RandomForestClassifier(bootstrap=True, class_weight=None, criterion='gini',
                                max_depth=None, max_features='auto', max_leaf_nodes=None,
                                min_impurity_decrease=0.0, min_impurity_split=None,
                                min_samples_leaf=1, min_samples_split=2,
                                min_weight_fraction_leaf=0.0, n_estimators=10, n_jobs=-1,
                                oob_score=False, random_state=0, verbose=0, warm_start=False)

['normal.' 'normal.' 'normal.' ..., 'normal.' 'normal.' 'normal.']
```

```
Out[21]: RandomForestClassifier(bootstrap=True, class_weight=None, criterion='gini',
                                max_depth=None, max_features='auto', max_leaf_nodes=None,
                                min_impurity_decrease=0.0, min_impurity_split=None,
                                min_samples_leaf=1, min_samples_split=2,
                                min_weight_fraction_leaf=0.0, n_estimators=10, n_jobs=-1,
                                oob_score=False, random_state=0, verbose=0, warm_start=False)

['normal.' 'normal.' 'normal.' ..., 'normal.' 'normal.' 'normal.']
```

```
In [22]: #Evaluation Metrics for Random Forest -- 2 labels
         #Confusion Matrix, accuracy, recall, precision and f1-score
         print(metrics.confusion_matrix(attack_labs_test, rdpreds))
         print(metrics.accuracy_score(attack_labs_test, rdpreds))
         print(metrics.recall_score(attack_labs_test, rdpreds,average='macro'))
         print(metrics.precision_score(attack_labs_test, rdpreds,average='macro'))
         print(metrics.f1_score(attack_labs_test, rdpreds,average='macro'))

         ht = sns.heatmap(pd.DataFrame(metrics.confusion_matrix(attack_labs_test, rdpreds)), ann
         ht.set(xlabel='Predicted Values', ylabel='True Values')
         plt.show()

         #Another way to view these metrics for every class label
         precision, recall, fscore, support = metrics.precision_recall_fscore_support(attack_lab

         print('precision: {}'.format(precision))
         print('recall: {}'.format(recall))
         print('fscore: {}'.format(fscore))
         print('support: {}'.format(support))
```

```

#Evaluation Metrics for Random Forest -- multi-labels
#Confusion Matrix, accuracy, recall, precision and f1-score
print(metrics.confusion_matrix(attack_labs_tall, rdpreds1))
print(metrics.accuracy_score(attack_labs_tall, rdpreds1))
print(metrics.recall_score(attack_labs_tall, rdpreds1,average='macro'))
print(metrics.precision_score(attack_labs_tall, rdpreds1,average='macro'))
print(metrics.f1_score(attack_labs_tall, rdpreds1,average='macro'))

ht = sns.heatmap(pd.DataFrame(metrics.confusion_matrix(attack_labs_tall, rdpreds1)), an
ht.set(xlabel='Predicted Values', ylabel='True Values')
plt.show()

#Another way to view these metrics for every class label
precision, recall, fscore, support = metrics.precision_recall_fscore_support(attack_lab

print('precision: {}'.format(precision))
print('recall: {}'.format(recall))
print('fscore: {}'.format(fscore))
print('support: {}'.format(support))

```

```


[[227882  22554]
 [   325 60268]]
0.926440942806
0.952288703459
0.863128452653
0.896335123835

```

```

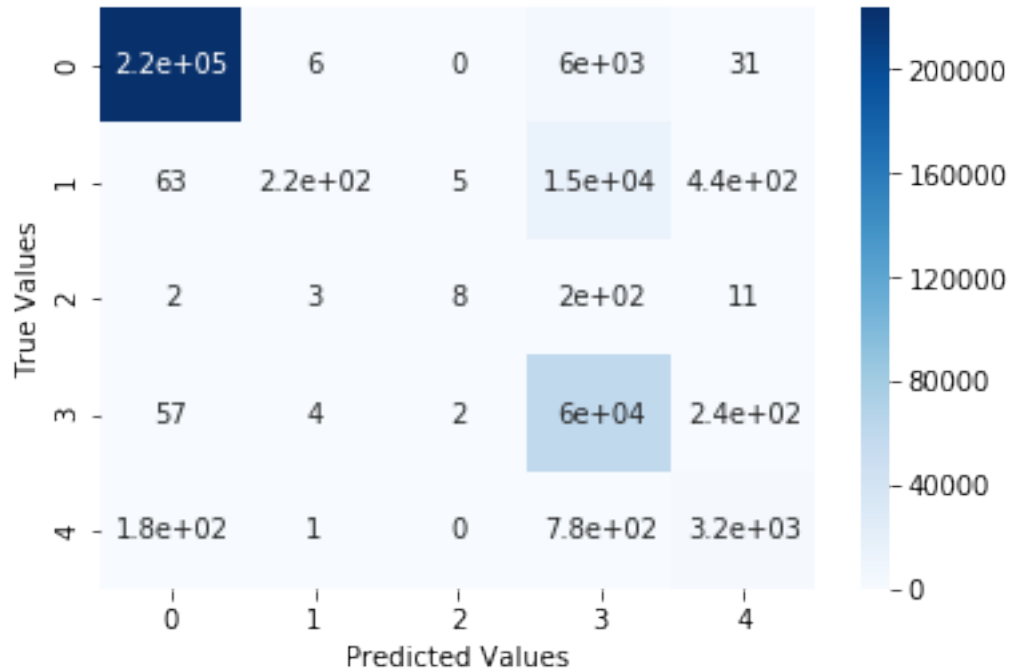
Out[22]: [Text(33,0.5,'True Values'), Text(0.5,15,'Predicted Values')]

```



```
precision: [ 0.99857585  0.72768105]
recall: [ 0.90994106  0.99463634]
fscore: [ 0.95220028  0.84046996]
support: [250436  60593]
[[223802      6      0  6014      31]
 [    63    223      5 15463     435]
 [     2      3      8   204      11]
 [    57      4      2 60295     235]
 [   178      1      0   781    3206]]
0.924460420089
0.557436410762
0.803955281697
0.542680957055
```

```
Out[22]: [Text(33,0.5,'True Values'), Text(0.5,15,'Predicted Values')]
```

```
precision: [ 0.99866132  0.94092827  0.53333333  0.72857885  0.81827463]
recall:    [ 0.97367448  0.01377479  0.03508772  0.99508194  0.76956313]
fscore:    [ 0.98600963  0.02715208  0.06584362  0.84122776  0.7931717 ]
support:   [229853  16189    228  60593   4166]
```

We see that the accuracy (92.64 percent) and F1 score (89.63 percent) for the 2-label classification has improved a lot as compared to the previous models. The F1 score for "normal" labels is also significantly higher (84.05 percent). However, in the multi-class classification, though the overall accuracy is higher, the recall values and F1 scores for "U2R" and "R2L" are low, leading to a lower F1 score overall.

1.1.9 Neural Networks

We now use a Neural Network model (Multi-Layer Perceptron) trained on the training dataset to predict "normal" and "bad" labels on the test dataset.

```
In [23]: #Multi Layer Perceptron Model -- 2 labels
mlpmod = MLPClassifier(activation='tanh',random_state=0)
mlpmod.fit(kdd_train_10per_scaled, attack_labs)
mlppreds = mlpmod.predict(kdd_test_10per_scaled)
print(mlppreds)

#Multi Layer Perceptron Model -- multi-labels
mlpmod1 = MLPClassifier(activation='tanh',random_state=0)
```

```
mlpmod1.fit(kdd_train_10per_scaled, attack_labs_all)
mlppreds1 = mlpmod1.predict(kdd_test_10per_scaled)
print(mlppreds1)
```

```
Out[23]: MLPClassifier(activation='tanh', alpha=0.0001, batch_size='auto', beta_1=0.9,
    beta_2=0.999, early_stopping=False, epsilon=1e-08,
    hidden_layer_sizes=(100,), learning_rate='constant',
    learning_rate_init=0.001, max_iter=200, momentum=0.9,
    nesterovs_momentum=True, power_t=0.5, random_state=0, shuffle=True,
    solver='adam', tol=0.0001, validation_fraction=0.1, verbose=False,
    warm_start=False)
```

```
['normal.' 'normal.' 'normal.' ..., 'normal.' 'normal.' 'normal.']
```

```
Out[23]: MLPClassifier(activation='tanh', alpha=0.0001, batch_size='auto', beta_1=0.9,
    beta_2=0.999, early_stopping=False, epsilon=1e-08,
    hidden_layer_sizes=(100,), learning_rate='constant',
    learning_rate_init=0.001, max_iter=200, momentum=0.9,
    nesterovs_momentum=True, power_t=0.5, random_state=0, shuffle=True,
    solver='adam', tol=0.0001, validation_fraction=0.1, verbose=False,
    warm_start=False)
```

```
['normal.' 'normal.' 'normal.' ..., 'normal.' 'normal.' 'normal.']
```

```
In [24]: #Evaluation Metrics for an MLP Classifier -- 2 labels
    #Confusion Matrix, accuracy, recall, precision and f1-score
    print(metrics.confusion_matrix(attack_labs_test, mlppreds))
    print(metrics.accuracy_score(attack_labs_test, mlppreds))
    print(metrics.recall_score(attack_labs_test, mlppreds, average='macro'))
    print(metrics.precision_score(attack_labs_test, mlppreds, average='macro'))
    print(metrics.f1_score(attack_labs_test, mlppreds, average='macro'))

    ht = sns.heatmap(pd.DataFrame(metrics.confusion_matrix(attack_labs_test, mlppreds)),
    ht.set(xlabel='Predicted Values', ylabel='True Values')
    plt.show()

    #Another way to view these metrics for every class label
    precision, recall, fscore, support = metrics.precision_recall_fscore_support(attack_labs_test, mlppreds)

    print('precision: {}'.format(precision))
    print('recall: {}'.format(recall))
    print('fscore: {}'.format(fscore))
    print('support: {}'.format(support))

    #Evaluation Metrics for an MLP Classifier -- multi-labels
    #Confusion Matrix, accuracy, recall, precision and f1-score
    print(metrics.confusion_matrix(attack_labs_tall, mlppreds1))
```

```

print(metrics.accuracy_score(attack_labs_tall, mlppreds1))
print(metrics.recall_score(attack_labs_tall, mlppreds1,average='macro'))
print(metrics.precision_score(attack_labs_tall, mlppreds1,average='macro'))
print(metrics.f1_score(attack_labs_tall, mlppreds1,average='macro'))

ht = sns.heatmap(pd.DataFrame(metrics.confusion_matrix(attack_labs_tall, mlppreds1)), a
ht.set(xlabel='Predicted Values', ylabel='True Values')
plt.show()

#Another way to view these metrics for every class label
precision, recall, fscore, support = metrics.precision_recall_fscore_support(attack_lab

print('precision: {}'.format(precision))
print('recall: {}'.format(recall))
print('fscore: {}'.format(fscore))
print('support: {}'.format(support))


```

```

[[227791  22645]
 [   966 59627]]
0.924087464513
0.946817630451
0.860265825686
0.892729833571

```

Out[24]: [Text(33,0.5,'True Values'), Text(0.5,15,'Predicted Values')]

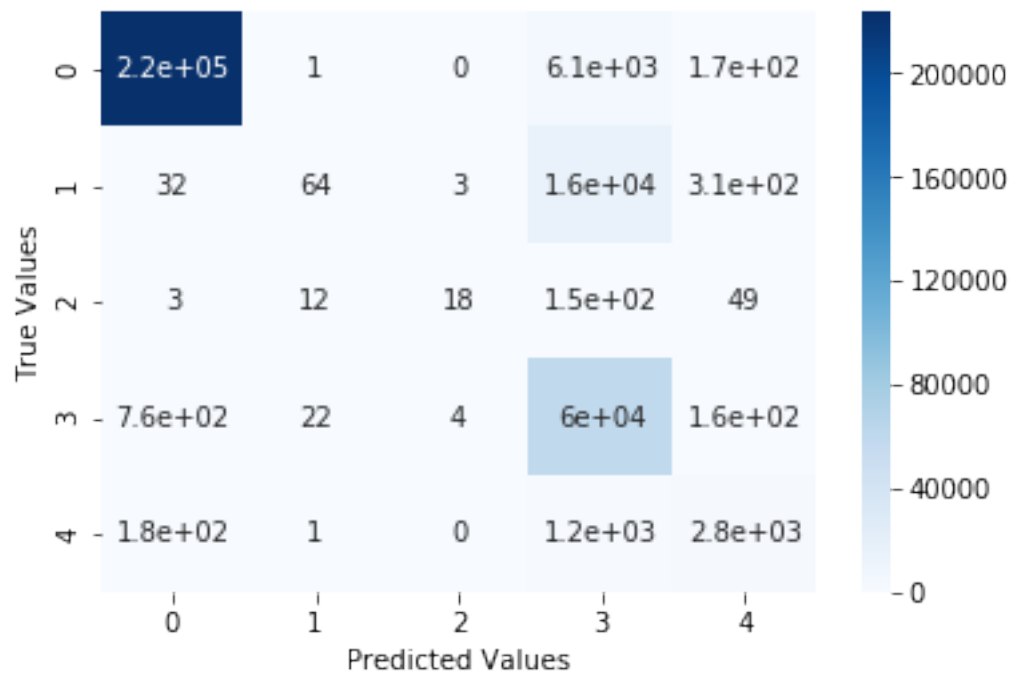


```

precision: [ 0.99577718  0.72475447]
recall: [ 0.9095777  0.98405756]
fscore: [ 0.95072758  0.83473209]
support: [250436  60593]
[[223577      1      0    6106    169]
 [    32     64      3   15778    312]
 [     3     12     18    146     49]
 [   758     22      4   59647    162]
 [   181      1      0    1195   2789]]
0.91983384186
0.541890202544
0.775323506598
0.539030370367

```

Out[24]: [Text(33,0.5,'True Values'), Text(0.5,15,'Predicted Values')]



```

precision: [ 0.99566246  0.64          0.72          0.71974853  0.80120655]
recall: [ 0.97269559  0.0039533  0.07894737  0.98438764  0.66946711]
fscore: [ 0.98404503  0.00785806  0.14229249  0.83151988  0.72943638]
support: [229853  16189    228  60593   4166]

```

We observe that the accuracy and F1 scores for a Multi-Layer Perceptron are very close to those for the Random Forest model for the 2-label as well as multi-label classification. The best

performance that we observe from among the implemented learning models (based on accuracy) is of the Random Forest model, followed closely by the Multi-Layer Perceptron model. The best performance (based on recall and F1 score) for multi-class classification is by the polynomial SVM model.

1.1.10 Ensemble Modeling

As we have already observed, Random Forest models provide the best results for the 2-label detection problem. However, for the multi-class detection problem, Polynomial SVM seem to provide a better F1 score. Here, we attempt to design an ensemble model of two random forests and one polynomial SVM model in order to improve the accuracy of the results.

```
In [27]: #Ensemble Modeling - Three different learning models
         #Test on 2 label as well as multi-class
```

```
rdforestmod = RandomForestClassifier(n_estimators=13,max_features='sqrt',n_jobs=-1, ran
rdforestmod1 = RandomForestClassifier(n_estimators=13,max_features='log2',n_jobs=-1, ra
polysvmmod = BaggingClassifier(SVC(kernel='poly',cache_size=7000),max_samples = 0.1, n_

modelest = []
modelest.append(('randomforest', rdforestmod))
modelest.append(('randomforestanother', rdforestmod1))
modelest.append(('polysvmmod', polysvmmod))

ensemblemod = VotingClassifier(modelest,voting='soft')

#2-label
ensemblemod.fit(kdd_train_10per_scaled, attack_labs)
enspreds = ensemblemod.predict(kdd_test_10per_scaled)
print(enspreds)

#multi-label
ensemblemod1 = VotingClassifier(modelest,voting='soft')
ensemblemod1.fit(kdd_train_10per_scaled, attack_labs_all)
enspreds1 = ensemblemod1.predict(kdd_test_10per_scaled)
print(enspreds1)
```

```
Out[27]: VotingClassifier(estimators=[('randomforest', RandomForestClassifier(bootstrap=True, cl
max_depth=None, max_features='sqrt', max_leaf_nodes=None,
min_impurity_decrease=0.0, min_impurity_split=None,
min_samples_leaf=1, min_samples_split=2,
...stimators=10, n_jobs=-1, oob_score=False,
random_state=None, verbose=0, warm_start=False))],
flatten_transform=None, n_jobs=1, voting='soft', weights=None)
```

```
['normal.' 'normal.' 'normal.' ..., 'normal.' 'normal.' 'normal.']
```

```
Out[27]: VotingClassifier(estimators=[('randomforest', RandomForestClassifier(bootstrap=True, cl
max_depth=None, max_features='sqrt', max_leaf_nodes=None,
```

```

        min_impurity_decrease=0.0, min_impurity_split=None,
        min_samples_leaf=1, min_samples_split=2,
        ...stimators=10, n_jobs=-1, oob_score=False,
        random_state=None, verbose=0, warm_start=False)),
        flatten_transform=None, n_jobs=1, voting='soft', weights=None)

['normal.' 'normal.' 'normal.' ..., 'normal.' 'normal.' 'normal.']

```

```

In [28]: #Evaluation Metrics for the ensemble model -- 2 labels
#Confusion Matrix, accuracy, recall, precision and f1-score
print(metrics.confusion_matrix(attack_labs_test, enspreds))
print(metrics.accuracy_score(attack_labs_test, enspreds))
print(metrics.recall_score(attack_labs_test, enspreds, average='macro'))
print(metrics.precision_score(attack_labs_test, enspreds, average='macro'))
print(metrics.f1_score(attack_labs_test, enspreds, average='macro'))

ht = sns.heatmap(pd.DataFrame(metrics.confusion_matrix(attack_labs_test, enspreds)),
ht.set(xlabel='Predicted Values', ylabel='True Values')
plt.show()

#Another way to view these metrics for every class label
precision, recall, fscore, support = metrics.precision_recall_fscore_support(attack_lab

print('precision: {}'.format(precision))
print('recall: {}'.format(recall))
print('fscore: {}'.format(fscore))
print('support: {}'.format(support))

#Evaluation Metrics for the ensemble model -- multi-labels
#Confusion Matrix, accuracy, recall, precision and f1-score
print(metrics.confusion_matrix(attack_labs_tall, enspreds1))
print(metrics.accuracy_score(attack_labs_tall, enspreds1))
print(metrics.recall_score(attack_labs_tall, enspreds1, average='macro'))
print(metrics.precision_score(attack_labs_tall, enspreds1, average='macro'))
print(metrics.f1_score(attack_labs_tall, enspreds1, average='macro'))

ht = sns.heatmap(pd.DataFrame(metrics.confusion_matrix(attack_labs_tall, enspreds1)), a
ht.set(xlabel='Predicted Values', ylabel='True Values')
plt.show()


#Another way to view these metrics for every class label
precision, recall, fscore, support = metrics.precision_recall_fscore_support(attack_lab

print('precision: {}'.format(precision))
print('recall: {}'.format(recall))
print('fscore: {}'.format(fscore))
print('support: {}'.format(support))

```


```
[[227909  22527]
 [   308 60285]]
0.926582408714
0.952482889678
0.863312306962
0.896529434911
```

Out[28]: [Text(33,0.5,'True Values'), Text(0.5,15,'Predicted Values')]



```
precision: [ 0.99865041  0.72797421]
recall: [ 0.91004887  0.9949169 ]
fscore: [ 0.95229321  0.84076566]
support: [250436  60593]
[[223496      0      0  6273    84]
 [      3    223      4 15626   333]
 [      2      5     12   182    27]
 [    508      0      1 59845   239]
 [    162      0      0   603  3401]]
0.922669590295
0.568555103316
0.847768494849
0.554080472589
```

```
Out[28]: [Text(33,0.5,'True Values'), Text(0.5,15,'Predicted Values')]
```



```
precision: [ 0.99698891  0.97807018  0.70588235  0.72513904  0.832762 ]
recall: [ 0.97234319  0.01377479  0.05263158  0.98765534  0.81637062]
fscore: [ 0.98451183  0.02716696  0.09795918  0.83627954  0.82448485]
support: [229853  16189    228  60593   4166]
```

We observe that the accuracy and F1 score have improved significantly for the 2-label classification over the previous models by using an ensemble model system. The F1 scores for multi-class classification too are higher than the Random Forest classifier.

1.1.11 Unsupervised Learning - Clustering

The above learning methods trained a model on the training data and using the training labels, which was used later to predict the corresponding labels for the test data. Here, we try a different learning technique where we use a clustering algorithm on the test data to identify patterns which we can later verify using the already provided target labels.

```
In [74]: #Clustering implemented using K-Means on test dataset (n=5 clusters)
```

```
kmeans_mod = KMeans(n_clusters=5)
kmeans_mod.fit(kdd_test_10per_scaled)
labels = kmeans_mod.labels_
```



```
Out[74]: KMeans(algorithm='auto', copy_x=True, init='k-means++', max_iter=300,
               n_clusters=5, n_init=10, n_jobs=1, precompute_distances='auto',
               random_state=None, tol=0.0001, verbose=0)
```

```
In [75]: pd.Series(labels).value_counts()
```

```
Out[75]: 3    164804
         0    101311
         1    44191
         2      722
         4        1
         dtype: int64
```

```
In [144]: #Check with target labels for the test dataset
```

```
cluster_labels = pd.DataFrame(labels)
clust_map = pd.concat([cluster_labels, attack_labs_test], axis=1)
clust_map.columns = ['cluster', 'attack_type']
type(clust_map)
clust_map.head(10)

df = clust_map.groupby(['cluster', 'attack_type']).size()
df1 = pd.DataFrame(df).reset_index()
df1.columns = ['cluster', 'attack_type', 'counts']
df1.head(20)

df.unstack().plot(kind='bar', stacked=True, color=['red', 'blue'], grid=False)
plt.xlabel('Clusters')
plt.ylabel('Counts')
plt.title('Cluster Distribution')
plt.show()
```

```
Out[144]: pandas.core.frame.DataFrame
```

```
Out[144]:   cluster attack_type
         0         0    normal.
         1         0    normal.
         2         0    normal.
         3         0     bad.
         4         0     bad.
         5         0     bad.
         6         0    normal.
         7         0    normal.
         8         0     bad.
         9         0    normal.
```

```
Out[144]:   cluster attack_type counts
         0         0     bad.  41572
         1         0    normal.  59739
```

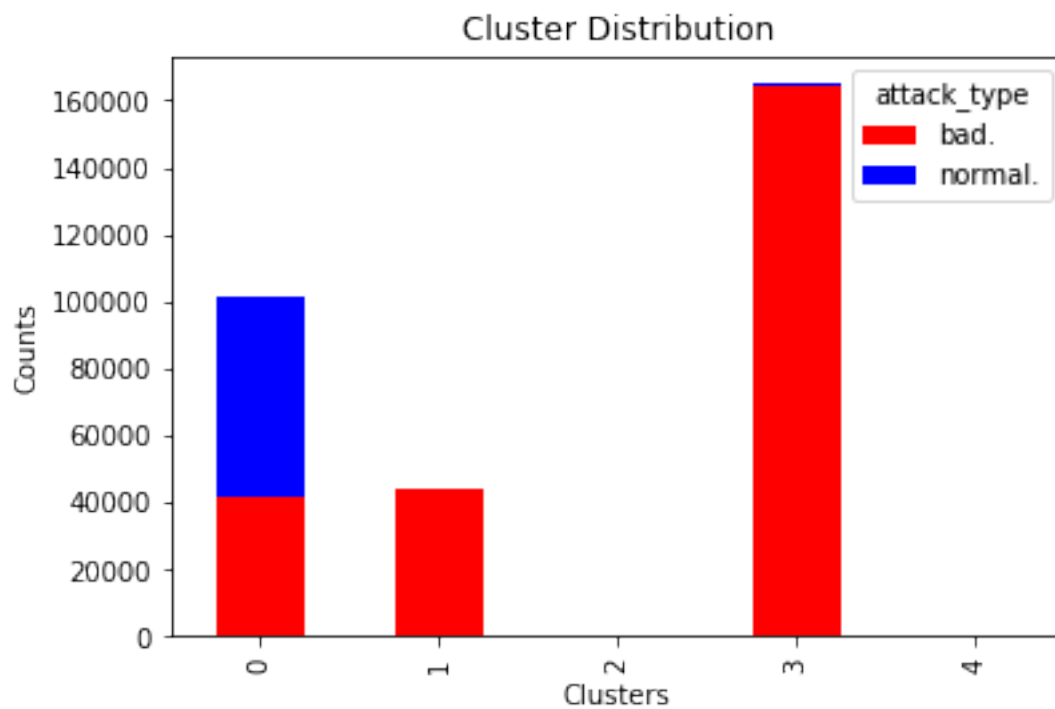
2	1	bad.	44076
3	1	normal.	115
4	2	bad.	717
5	2	normal.	5
6	3	bad.	164071
7	3	normal.	733
8	4	normal.	1

Out[144]: <matplotlib.axes._subplots.AxesSubplot at 0x266a648ca90>

Out[144]: Text(0.5,0,'Clusters')

Out[144]: Text(0,0.5,'Counts')

Out[144]: Text(0.5,1,'Cluster Distribution')



In [145]: *#Doing this for all major attack types in the test dataset*

```

clust_map = pd.concat([cluster_labels,attack_labs_tall],axis=1)
clust_map.columns = ['cluster','attack_type']
type(clust_map)
clust_map.head(10)

df = clust_map.groupby(['cluster','attack_type']).size()

```

```

df1 = pd.DataFrame(df).reset_index()
df1.columns = ['cluster', 'attack_type', 'counts']
df1.head(20)

df.unstack().plot(kind='bar', stacked=True, color=['red', 'blue', 'green', 'yellow', 'orange'])
plt.xlabel('Clusters')
plt.ylabel('Counts')
plt.title('Cluster Distribution')
plt.show()

```

Out[145]: pandas.core.frame.DataFrame

```

Out[145]:   cluster attack_type
0         0      normal.
1         0      normal.
2         0      normal.
3         0        R2L.
4         0        R2L.
5         0        R2L.
6         0      normal.
7         0      normal.
8         0        R2L.
9         0      normal.

```

```

Out[145]:   cluster attack_type counts
0         0        DOS.   24831
1         0        R2L.   15465
2         0        U2R     89
3         0      normal.   59739
4         0      probe.    1187
5         1        DOS.  40955
6         1        R2L.     9
7         1        U2R    136
8         1      normal.    115
9         1      probe.   2976
10        2        R2L.    714
11        2        U2R     3
12        2      normal.     5
13        3        DOS. 164067
14        3        R2L.     1
15        3      normal.    733
16        3      probe.     3
17        4      normal.     1

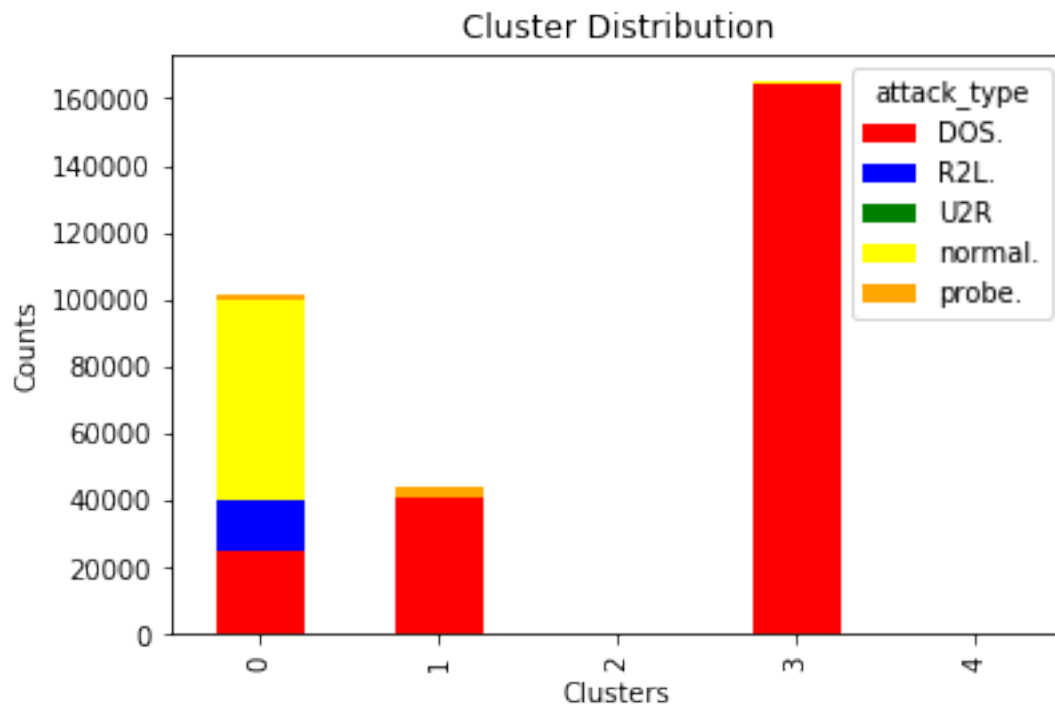
```

Out[145]: <matplotlib.axes._subplots.AxesSubplot at 0x26746543c88>

Out[145]: Text(0.5,0,'Clusters')

Out[145]: Text(0,0.5,'Counts')

```
Out[145]: Text(0.5,1,'Cluster Distribution')
```



We see that most of the clusters show one particular type of attack or normal as the dominant type in the cluster. We can do it for more clusters for greater granularity.

```
In [148]: #Clustering for n=10 clusters
```

```
kmeans_mod1 = KMeans(n_clusters=10)
kmeans_mod1.fit(kdd_test_10per_scaled)
labels1 = kmeans_mod1.labels_

#Check with target labels for the test dataset

cluster_labels1 = pd.DataFrame(labels1)
clust_map1 = pd.concat([cluster_labels1,attack_labs_tall],axis=1)
clust_map2 = pd.concat([cluster_labels1,attack_labs_test],axis=1)
clust_map1.columns = ['cluster','attack_type']
clust_map2.columns = ['cluster','attack_type']

df = clust_map2.groupby(['cluster','attack_type']).size()
df1 = pd.DataFrame(df).reset_index()
df1.columns = ['cluster','attack_type','counts']
df1.head(50)

df.unstack().plot(kind='bar',stacked=True, color=['red','blue'], grid=False)
```

```

plt.xlabel('Clusters')
plt.ylabel('Counts')
plt.title('Cluster Distribution')
plt.show()

df = clust_map1.groupby(['cluster', 'attack_type']).size()
df1 = pd.DataFrame(df).reset_index()
df1.columns = ['cluster', 'attack_type', 'counts']
df1.head(50)

df.unstack().plot(kind='bar', stacked=True, color=['red', 'blue', 'green', 'yellow', 'orange'])
plt.xlabel('Clusters')
plt.ylabel('Counts')
plt.title('Cluster Distribution')
plt.show()

```

```

Out[148]: KMeans(algorithm='auto', copy_x=True, init='k-means++', max_iter=300,
                n_clusters=10, n_init=10, n_jobs=1, precompute_distances='auto',
                random_state=None, tol=0.0001, verbose=0)

```

```

Out[148]:
   cluster attack_type  counts
0         0         bad. 164005
1         0      normal.    728
2         1         bad.  41470
3         1      normal.    95
4         2         bad.   717
5         2      normal.     5
6         3         bad.     5
7         4         bad. 18096
8         4      normal.    23
9         5      normal.     1
10        6         bad. 23043
11        6      normal. 59357
12        7         bad.  3081
13        7      normal.   380
14        8      normal.     3
15        9         bad.    19
16        9      normal.     1

```

```

Out[148]: <matplotlib.axes._subplots.AxesSubplot at 0x268bb5a8d30>

```

```

Out[148]: Text(0.5,0,'Clusters')

```

```


Out[148]: Text(0,0.5,'Counts')

```

```

Out[148]: Text(0.5,1,'Cluster Distribution')

```



```

Out[148]:
  cluster attack_type counts
0        0      DOS. 164001
1        0      R2L.      1
2        0   normal.    728
3        0   probe.      3
4        1      DOS. 40905
5        1      R2L.      5
6        1      U2R.   117
7        1   normal.    95
8        1   probe.   443
9        2      R2L.   714
10       2      U2R.      3
11       2   normal.      5
12       3      U2R.      5
13       4      DOS. 17813
14       4      R2L.      3
15       4   normal.    23
16       4   probe.   280
17       5   normal.      1
18       6      DOS.   713
19       6      R2L. 15387
20       6      U2R.    66
21       6   normal. 59357
22       6   probe.   477

```

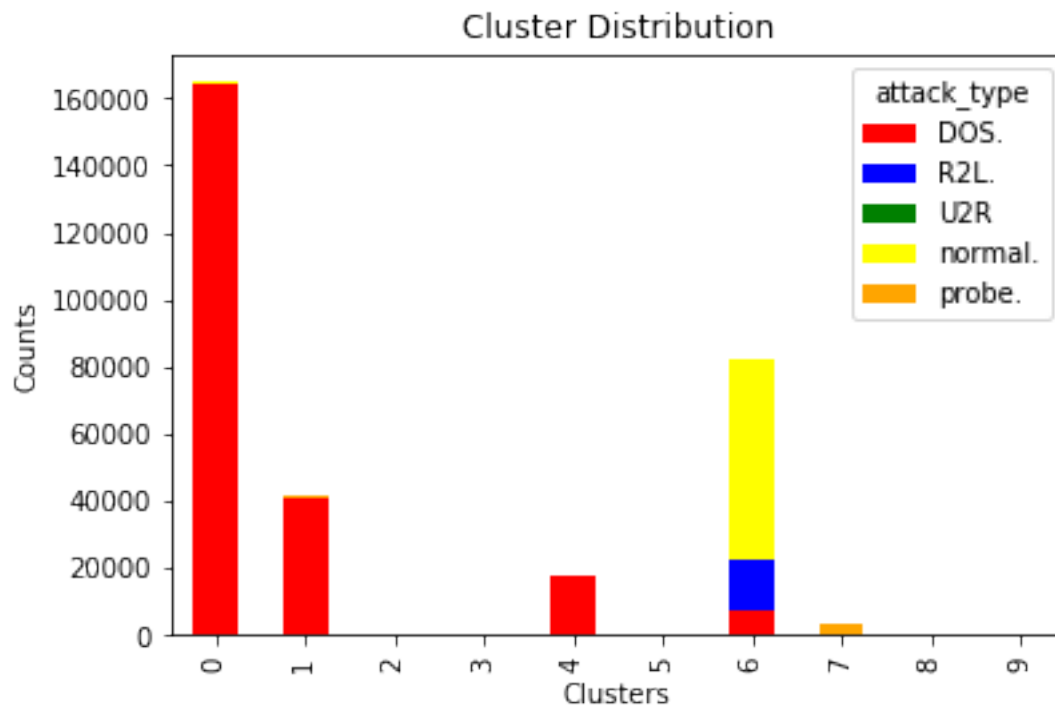
23	7	DOS.	21
24	7	R2L.	77
25	7	U2R	20
26	7	normal.	380
27	7	probe.	2963
28	8	normal.	3
29	9	R2L.	2
30	9	U2R	17
31	9	normal.	1

Out[148]: <matplotlib.axes._subplots.AxesSubplot at 0x267463c2198>

Out[148]: Text(0.5,0,'Clusters')

Out[148]: Text(0,0.5,'Counts')

Out[148]: Text(0.5,1,'Cluster Distribution')



We observe similar results here as well. Hence, we see that even if we do not have a training dataset to train a supervised model on, unsupervised algorithms like clustering can give us a fairly good idea regarding the classification of "bad" intrusions from "normal" traffic signatures.